



Machine perception

Derivatives and edge detection

Matej Kristan



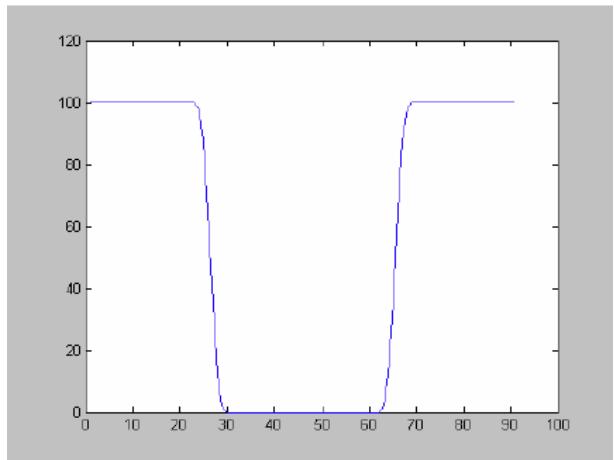
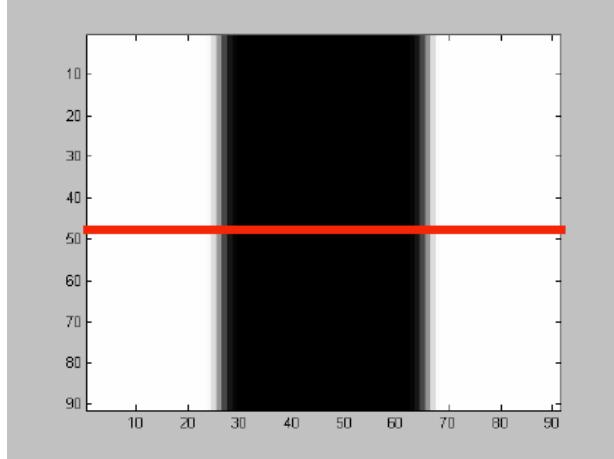
Laboratorij za Umetne Vizualne Spoznavne Sisteme,
Fakulteta za računalništvo in informatiko,
Univerza v Ljubljani



Machine perception

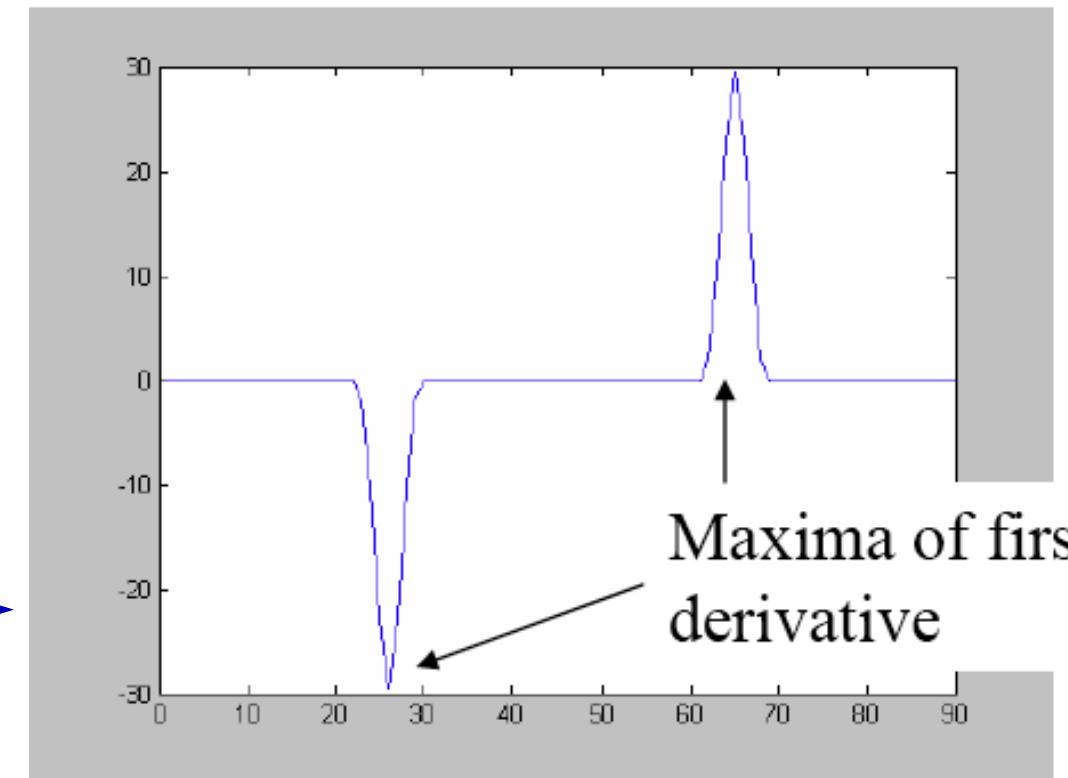
IMAGE DERIVATIVES

1D derivative: Intuition



First order derivative

$$\frac{d}{dx} f(x)$$



Maxima of first
derivative

Derivatives and convolution

- A partial derivative of a continuous 2D function $f(x,y)$:

$$\frac{\partial f(x, y)}{\partial x} = \lim_{\varepsilon \rightarrow 0} \frac{f(x + \varepsilon, y) - f(x, y)}{\varepsilon}$$

- For discrete case, approximate by using finite differences:

$$\frac{\partial f(x, y)}{\partial x} \approx \frac{f(x+1, y) - f(x, y)}{1}$$

- Question: If implemented by convolution, what would the convolution kernel for derivative look like? (Next slide)

Partial derivatives: Implementation



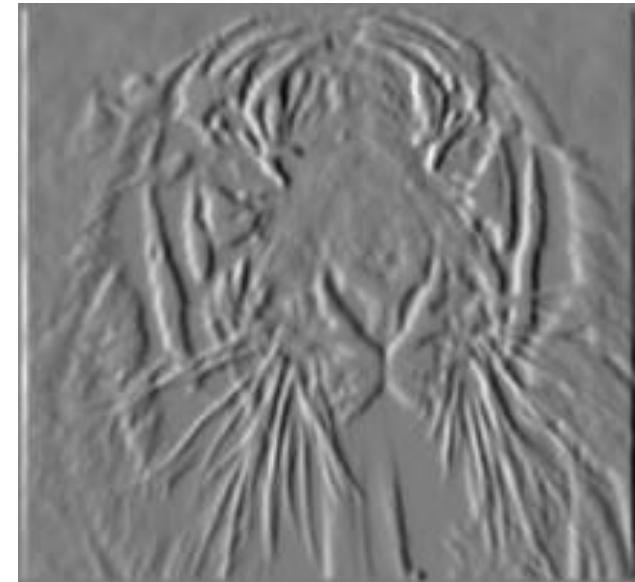
Horizontal derivative

$$\frac{\partial f(x,y)}{\partial x} \approx \frac{f(x+1,y) - f(x,y)}{1}$$

*

-1	1
----	---

 =



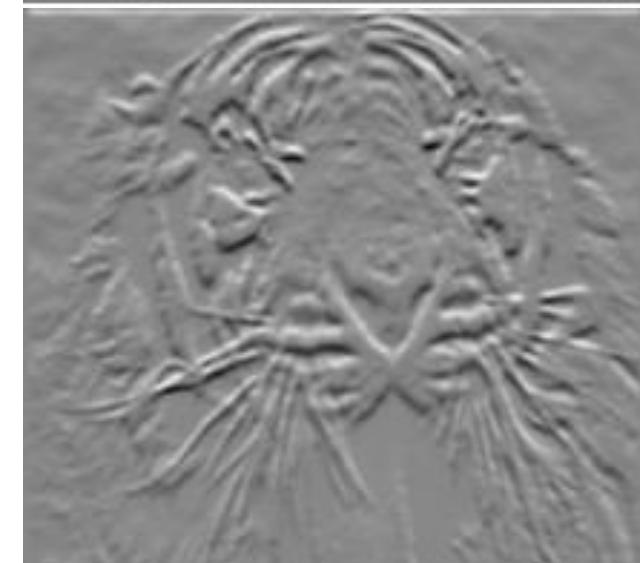
Vertical derivative

$$\frac{\partial f(x,y)}{\partial y} \approx \frac{f(x,y+1) - f(x,y)}{1}$$

*

-1
1

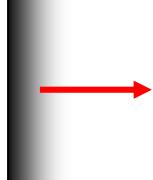
 =



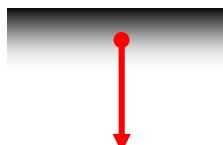
Partial derivatives: Image gradient

- Image gradient: $\nabla f = \left[\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right]$
- Gradient points in direction of greatest intensity change:

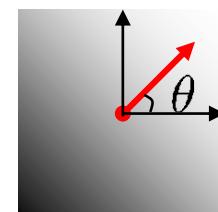
$$\nabla f = \left[\frac{\partial f}{\partial x}, 0 \right]$$



$$\nabla f = \left[0, \frac{\partial f}{\partial y} \right]$$



$$\nabla f = \left[\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right]$$



- Gradient direction (orientation of edge normal):

$$\theta = \tan^{-1} \left(\frac{\partial f}{\partial y} / \frac{\partial f}{\partial x} \right)$$

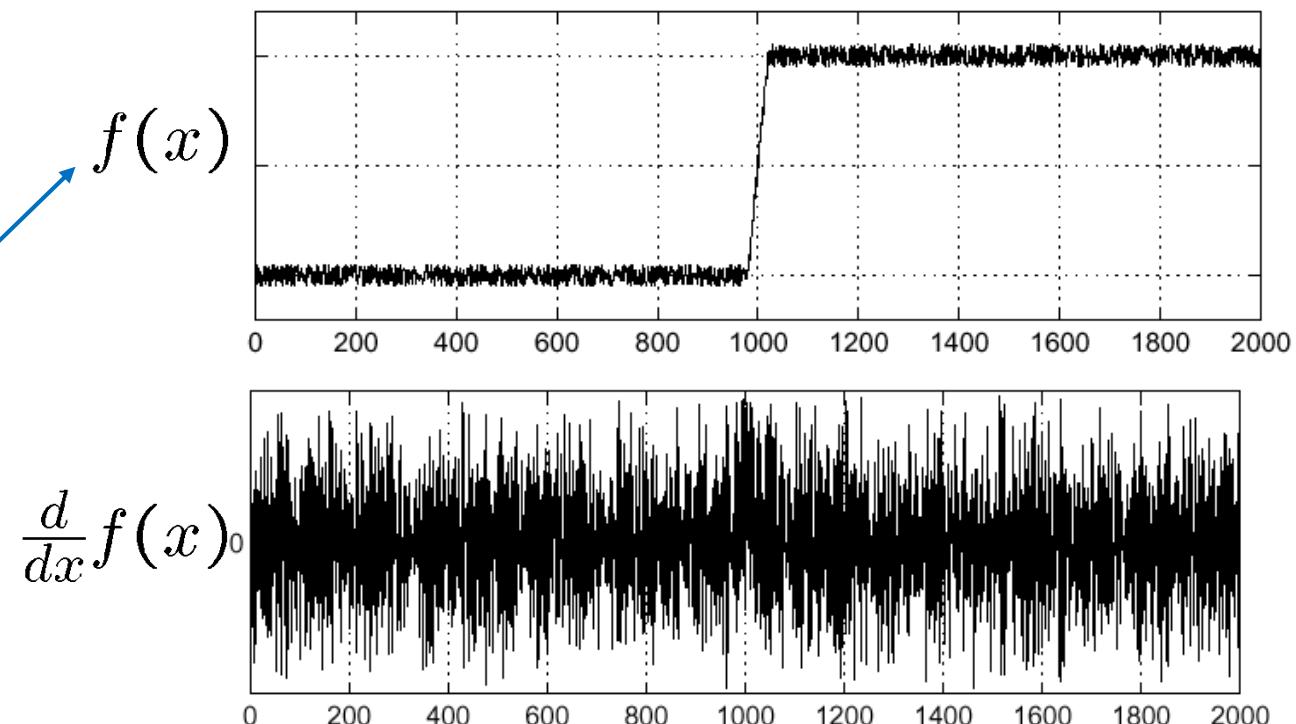
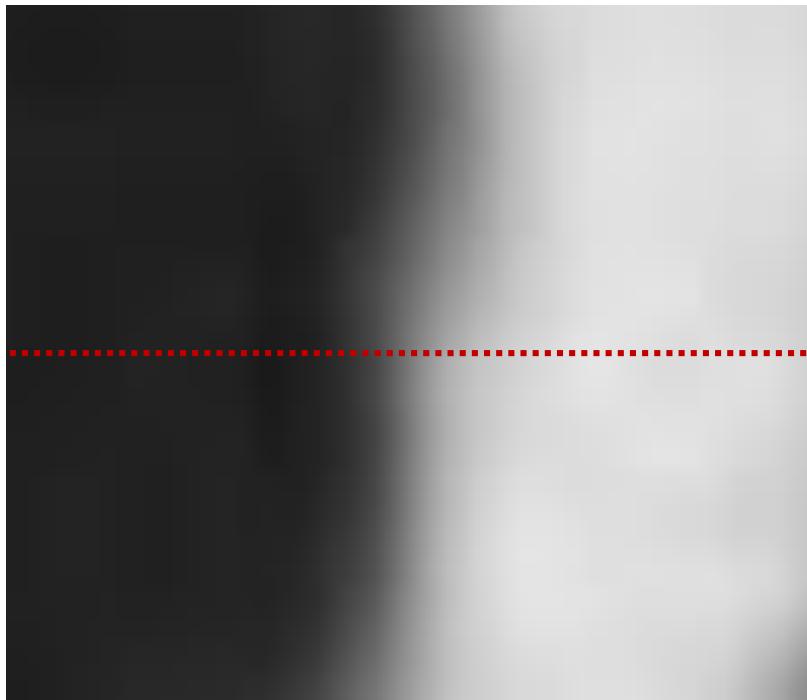
- Gradient strength is defined by its magnitude:

$$\|\nabla f\| = \sqrt{\left(\frac{\partial f}{\partial x} \right)^2 + \left(\frac{\partial f}{\partial y} \right)^2}$$



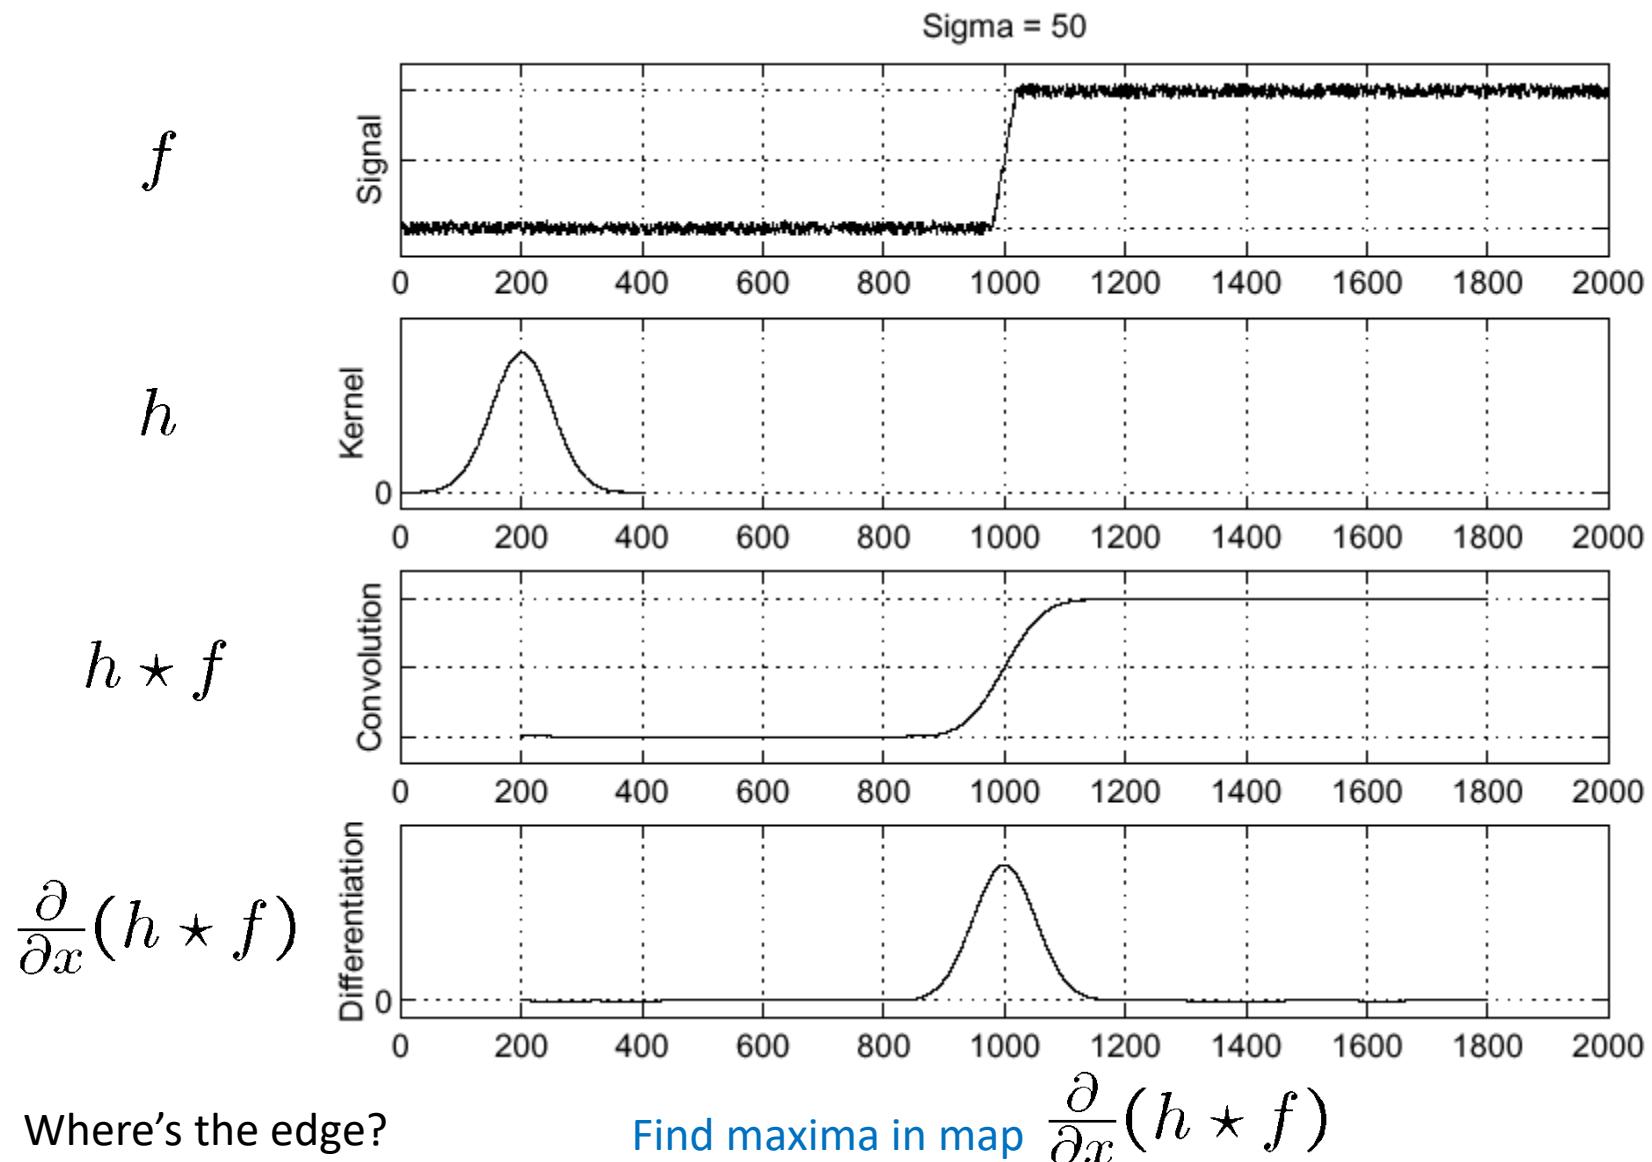
Discrete world is noisy...

- Take a **single line** in the image:
 - Plot **intensities** w.r.t. pixels:



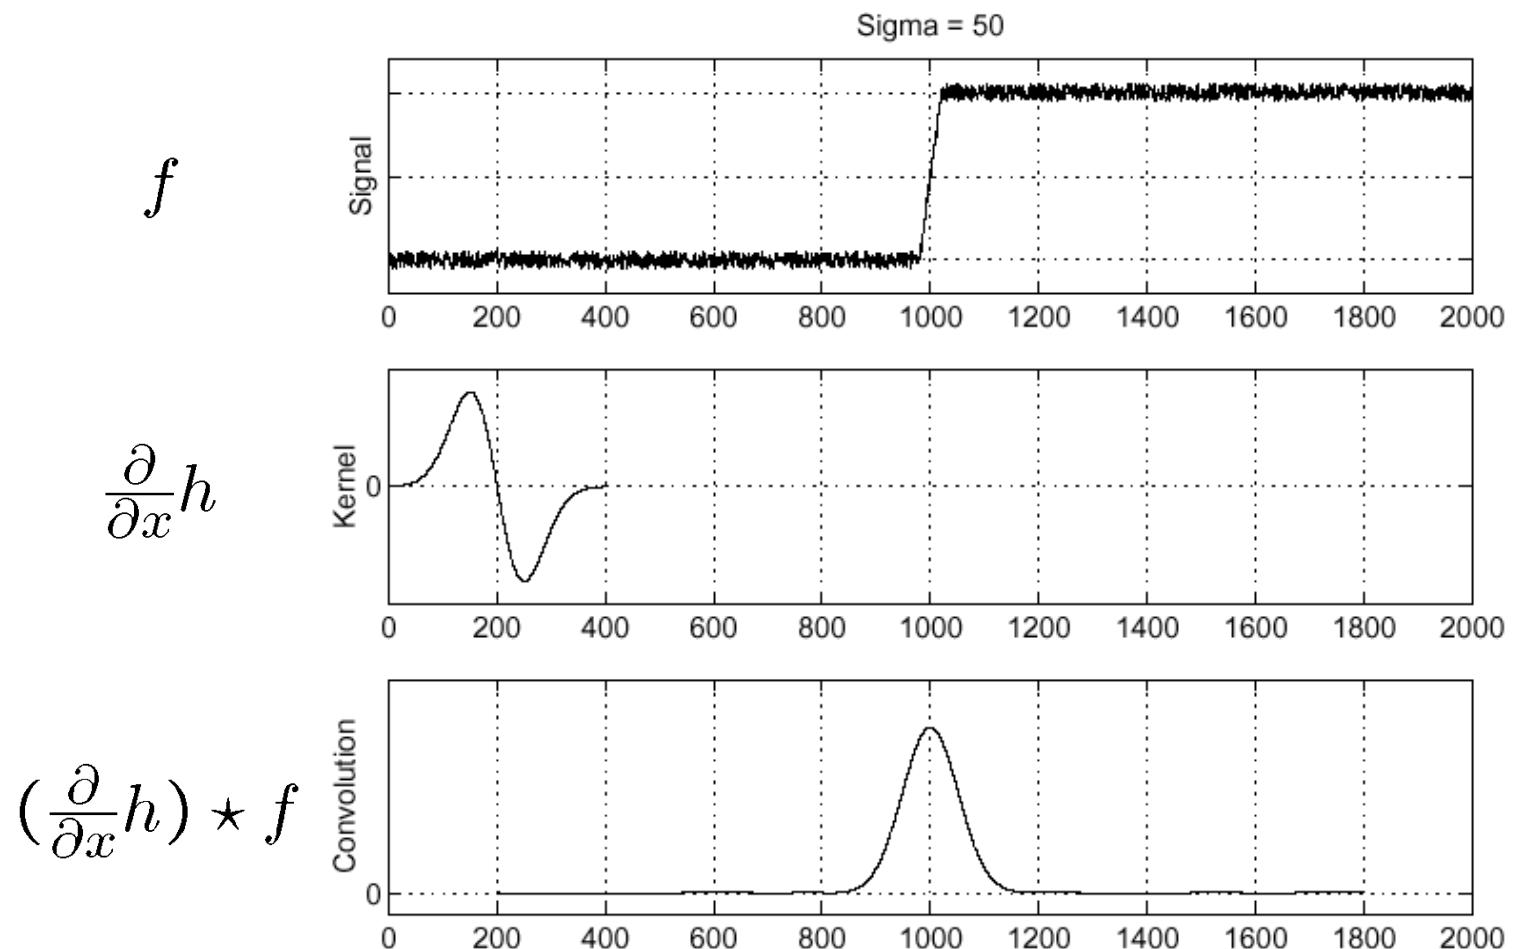
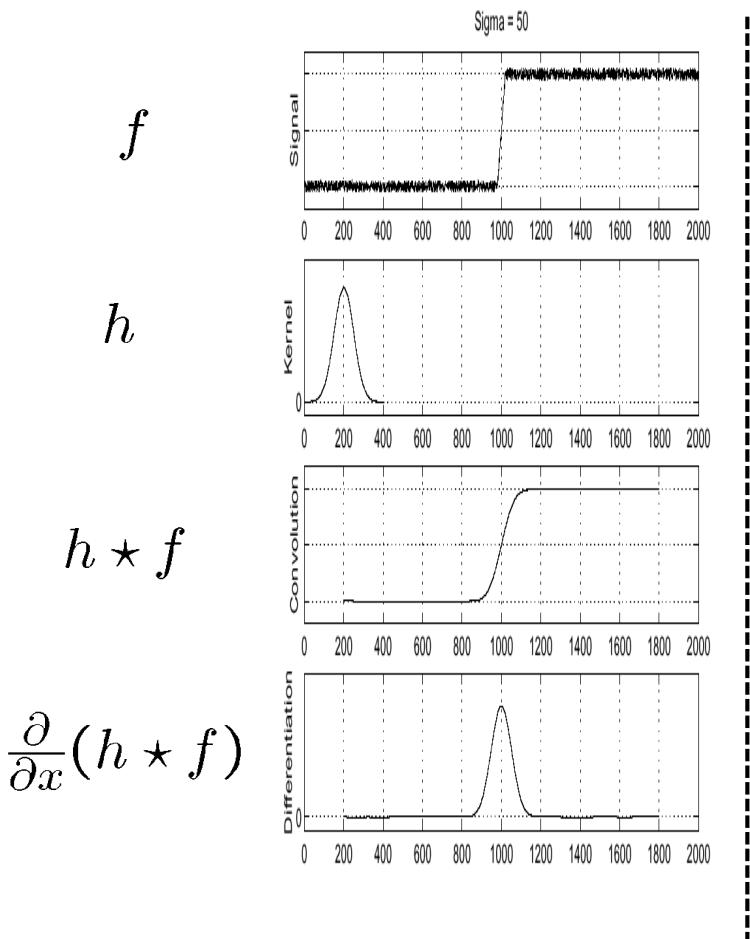
So where did the edge go!? Noise gets amplified by derivation...

Solution: Smooth the image first



Remember convolution properties

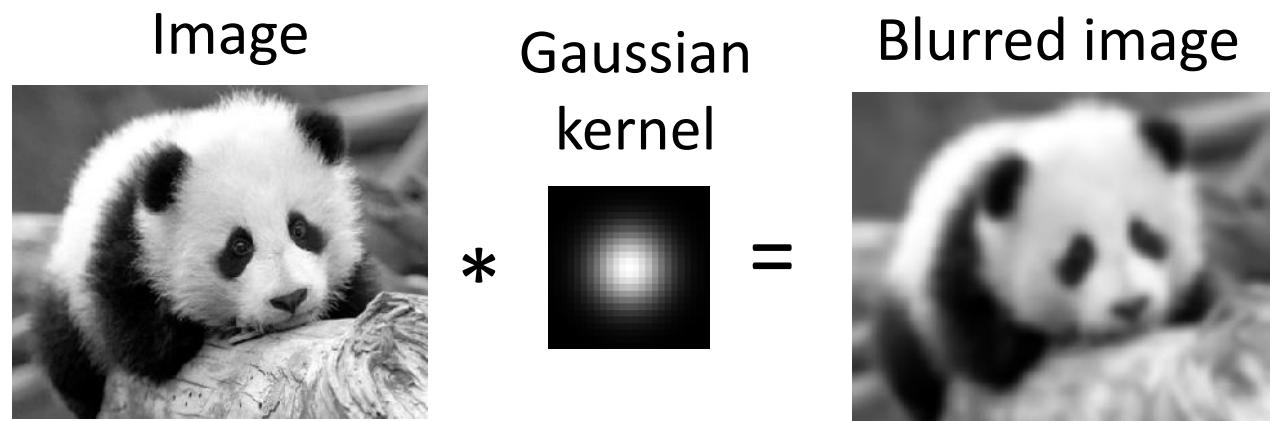
- Derivatives: $\frac{\partial}{\partial x}(h \star f) = (\frac{\partial}{\partial x}h) \star f$



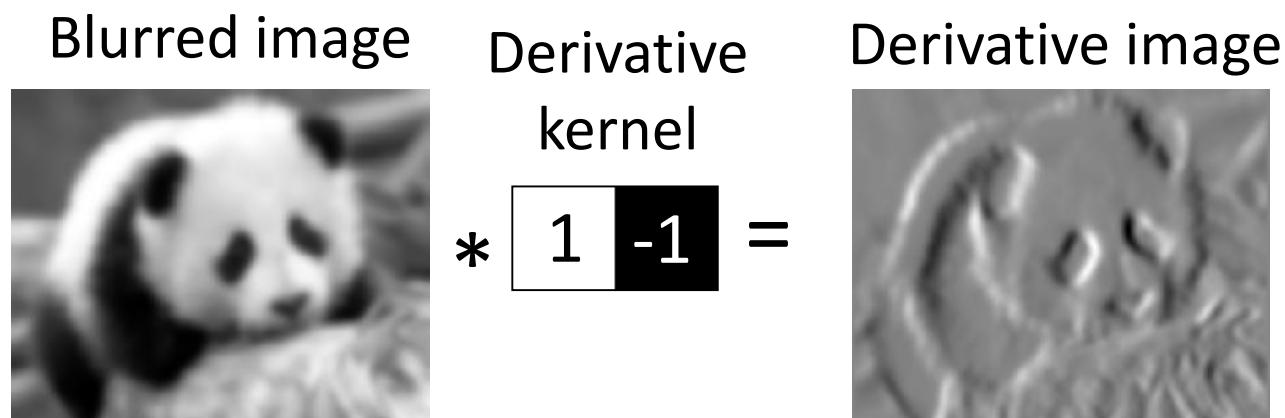
2D partial derivatives – naive approach

1. Smooth the image by a 2D Gaussian filter
2. Take derivative w.r.t. x

1 Blurring

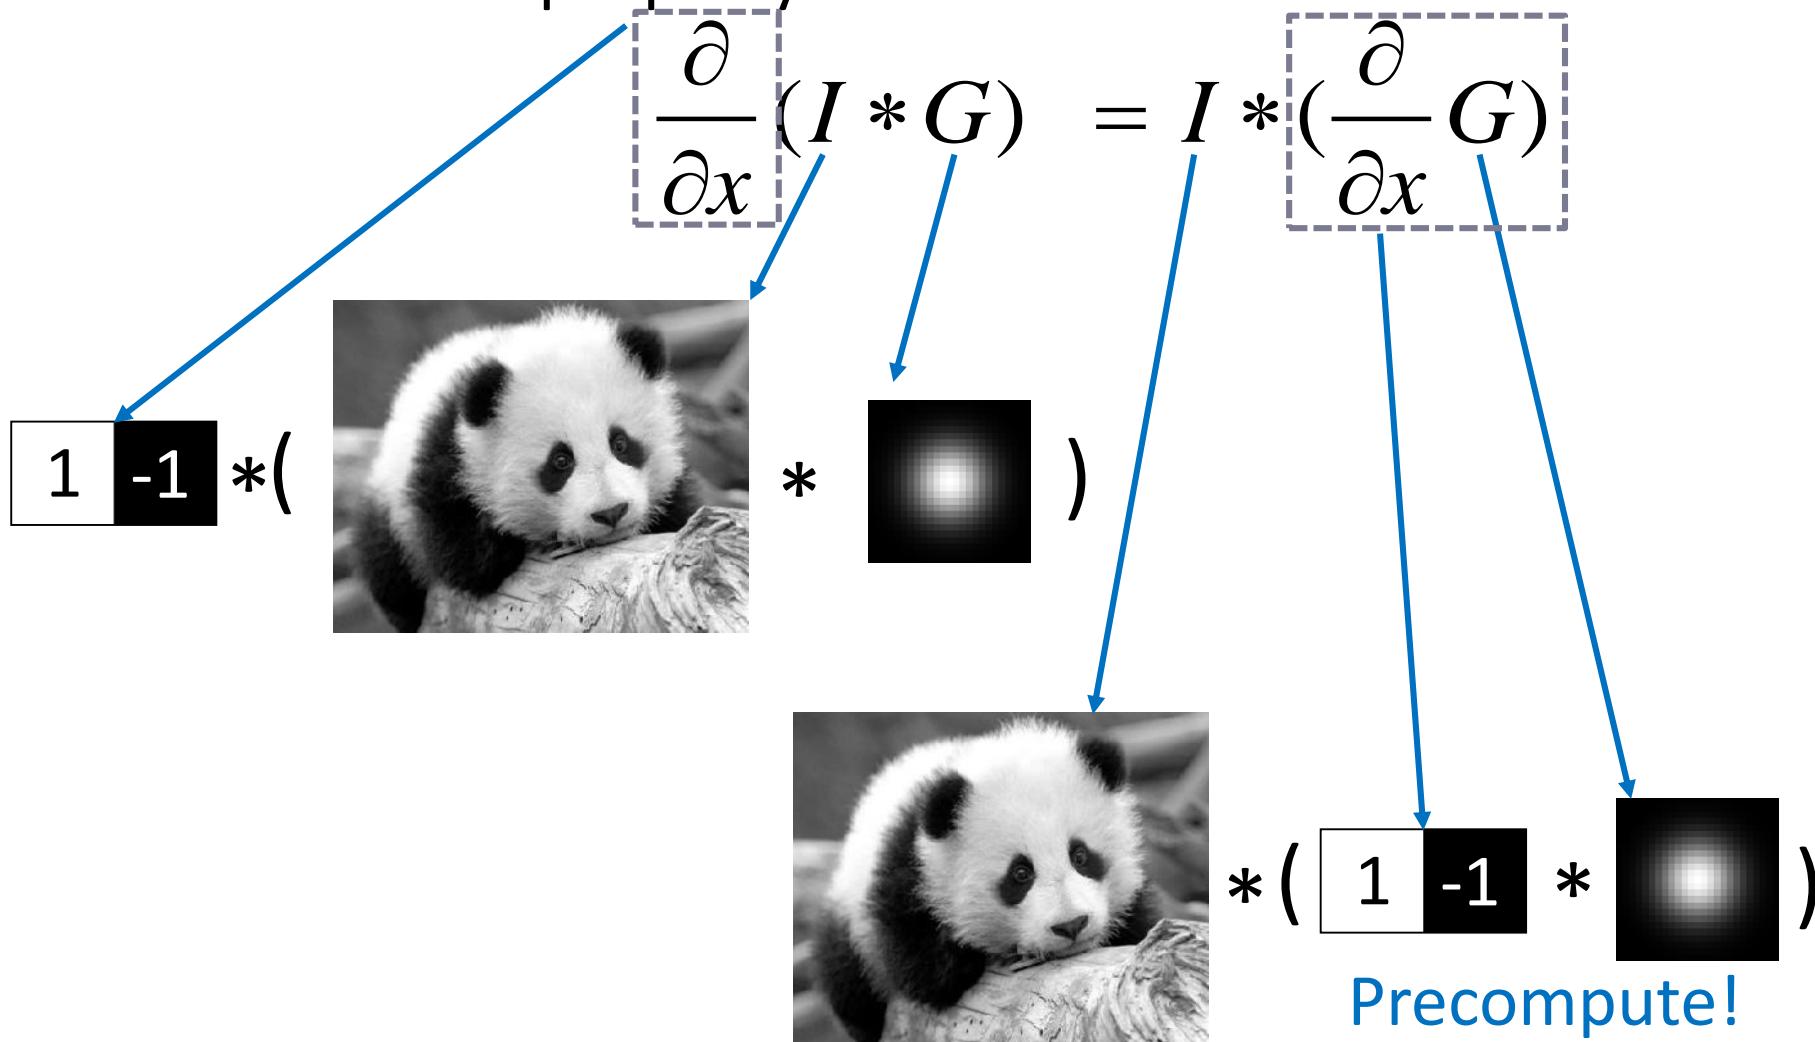


2 Differentiating
w.r.t. x



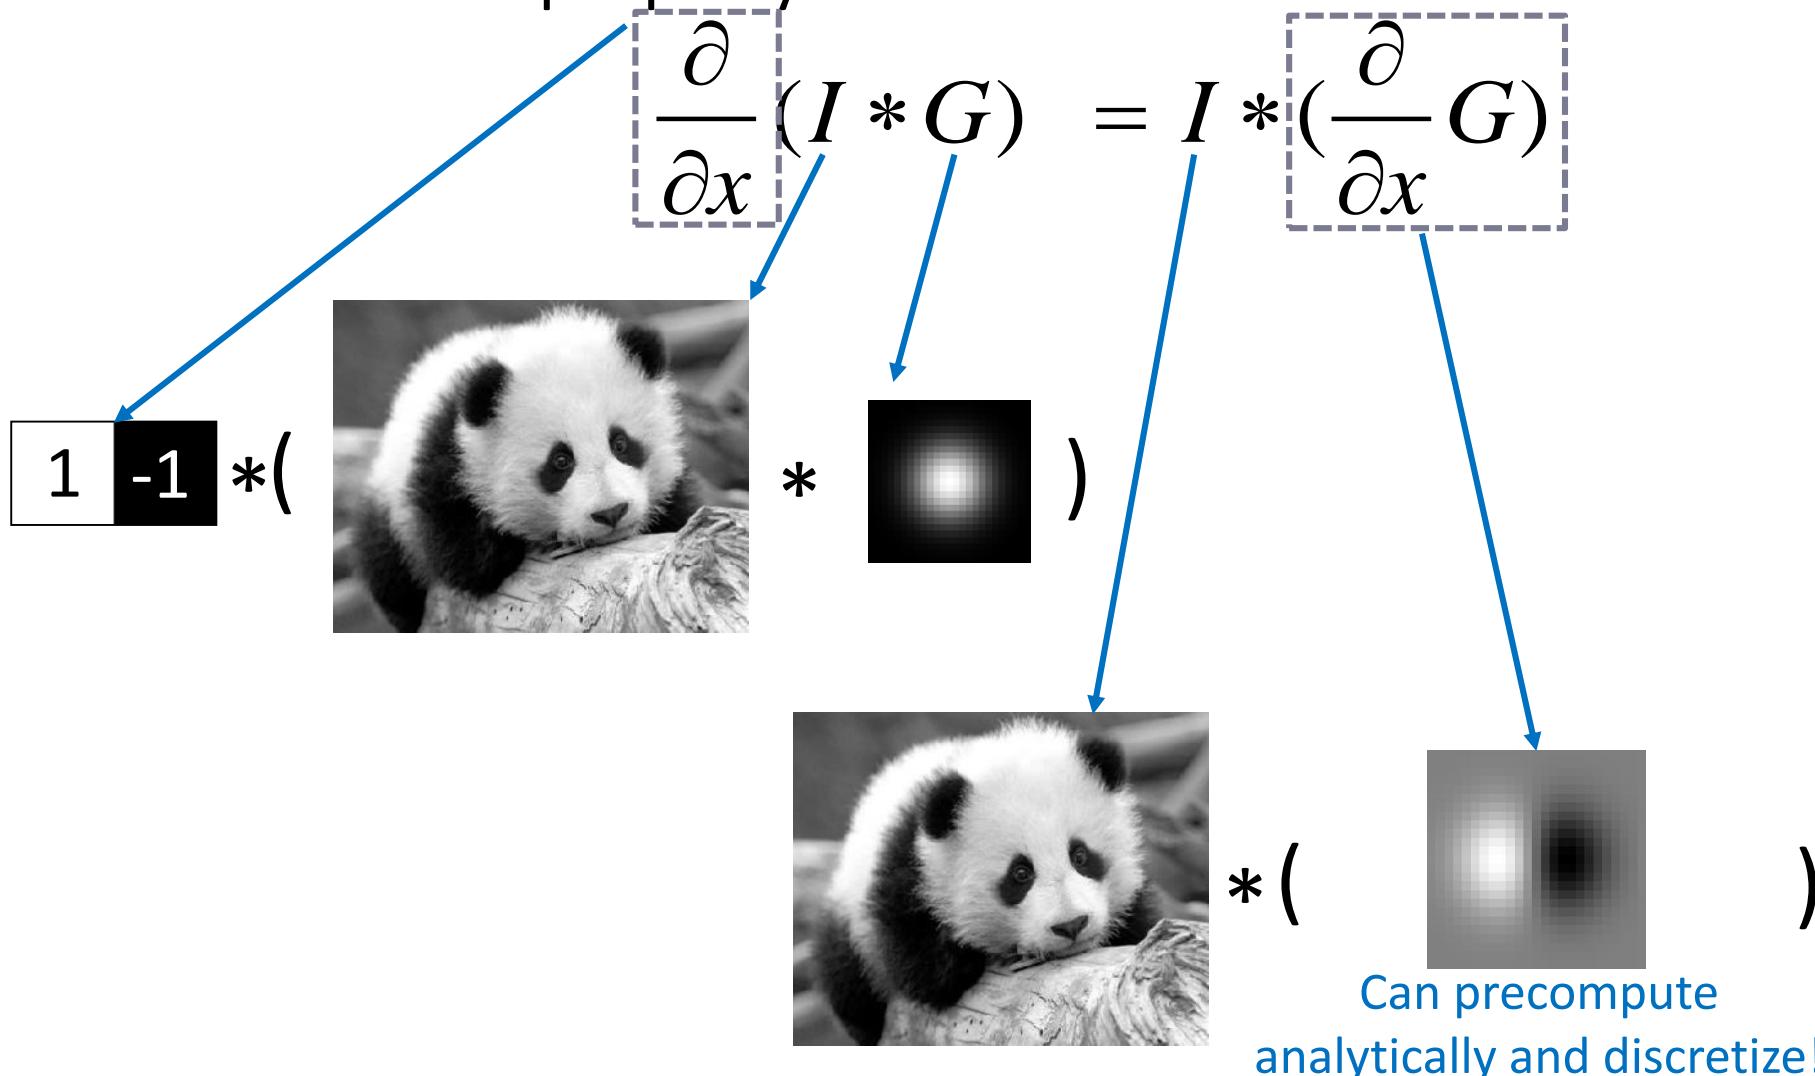
2D partial derivatives – smarter approach

- Recall the convolution property:



Smarter way

- Recall the convolution property:



2D partial derivatives – smarter approach

- Recall the convolution property:

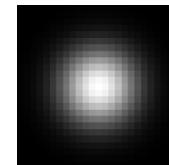
$$\frac{\partial}{\partial x}(I * G) = I * \left(\frac{\partial}{\partial x}G\right)$$

Naiive:

$$\begin{bmatrix} 1 & -1 \end{bmatrix} * ($$



*



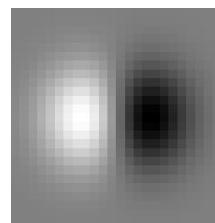
) =



Smarter:



*

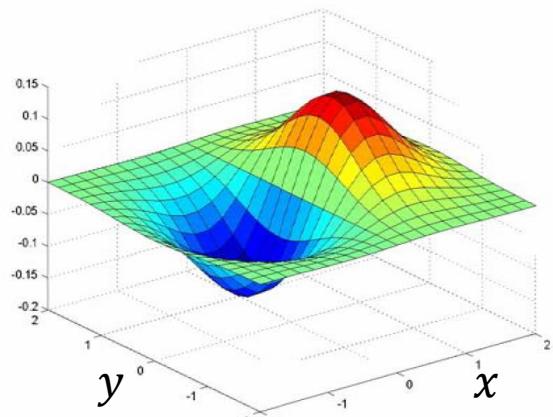


=

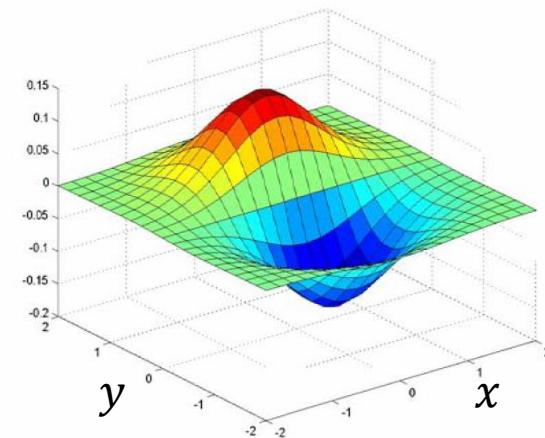


Gaussian partial derivatives

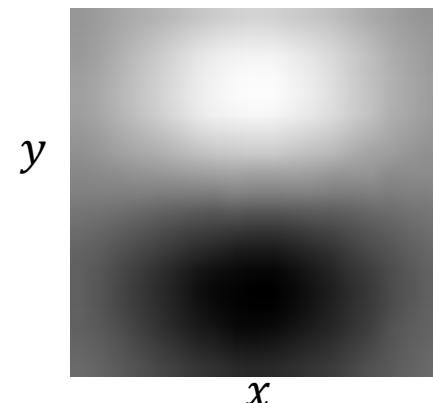
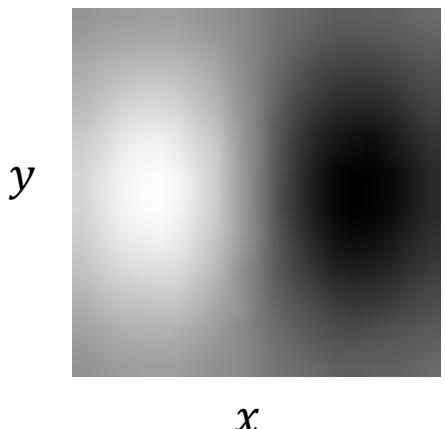
- Convolution kernels for taking partial derivatives w.r.t. x and y :



w.r.t. x



w.r.t. y



Some other popular kernels

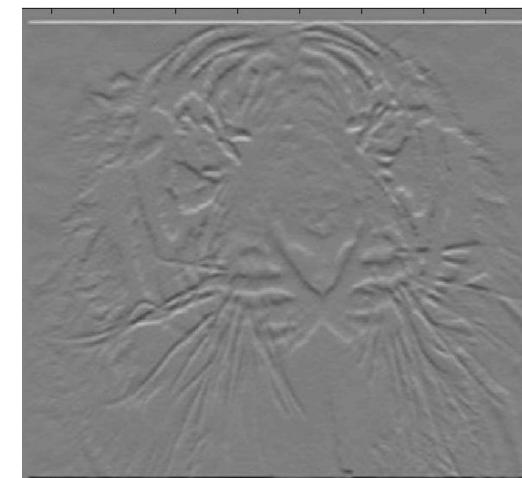
Prewitt: $M_x = \begin{bmatrix} -1 & 0 & 1 \\ -1 & 0 & 1 \\ -1 & 0 & 1 \end{bmatrix}$; $M_y = \begin{bmatrix} 1 & 1 & 1 \\ 0 & 0 & 0 \\ -1 & -1 & -1 \end{bmatrix}$

Sobel: $M_x = \begin{bmatrix} -1 & 0 & 1 \\ -2 & 0 & 2 \\ -1 & 0 & 1 \end{bmatrix}$; $M_y = \begin{bmatrix} 1 & 2 & 1 \\ 0 & 0 & 0 \\ -1 & -2 & -1 \end{bmatrix}$



Roberts: $M_x = \begin{bmatrix} 0 & 1 \\ -1 & 0 \end{bmatrix}$; $M_y = \begin{bmatrix} 1 & 0 \\ 0 & -1 \end{bmatrix}$

```
>> My = fspecial('sobel');  
>> outim = imfilter(double(im), My);  
>> imagesc(outim);  
>> colormap gray;
```



Laplacian of Gaussian (LoG)

- How does this function

h ... Gaussian

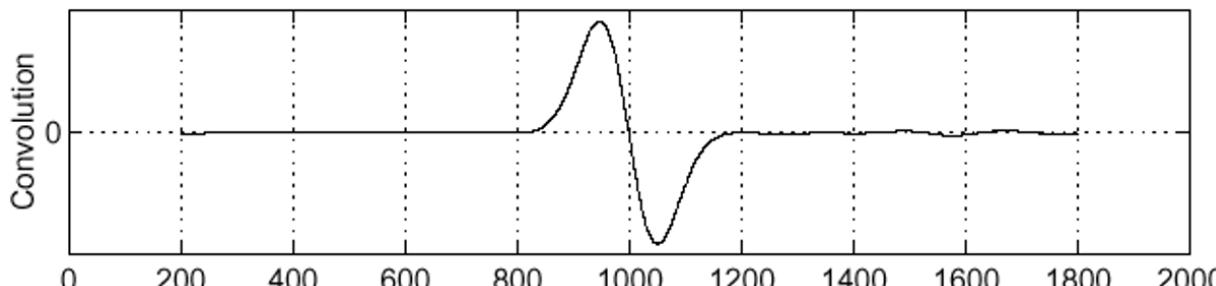
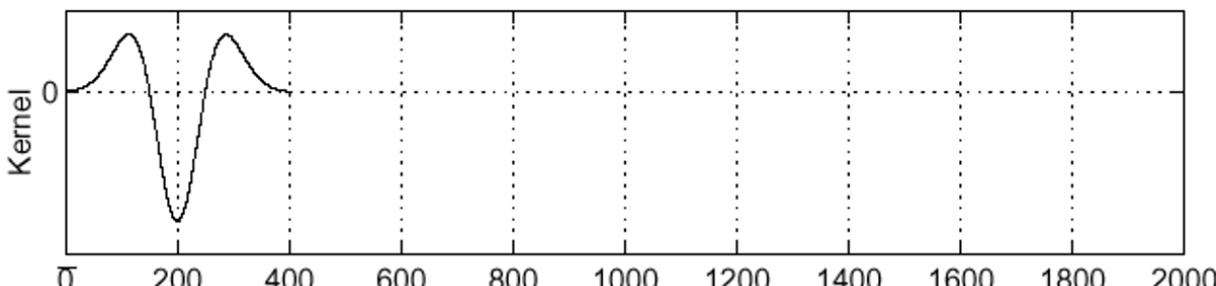
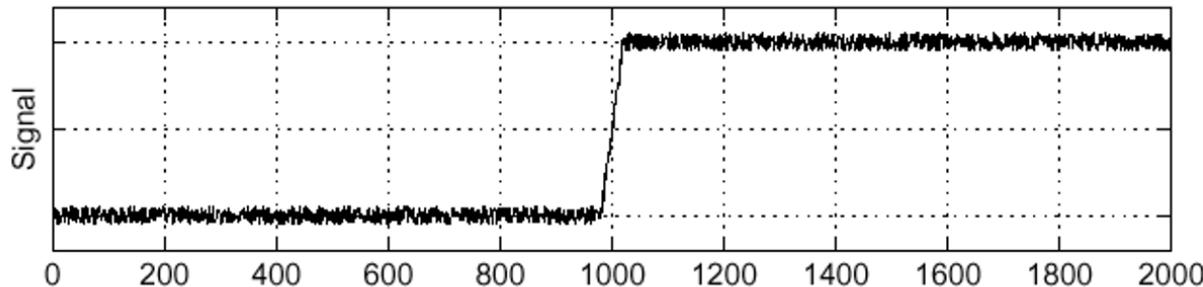
f

$$\frac{\partial^2}{\partial x^2} h$$

$$(\frac{\partial^2}{\partial x^2} h) \star f$$

$$\frac{\partial^2}{\partial x^2}(h \star f)?$$

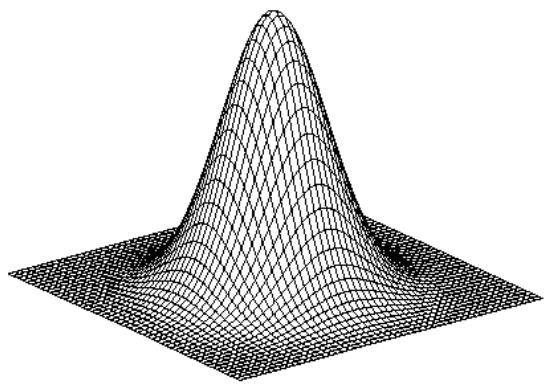
Sigma = 50



Where's the edge?

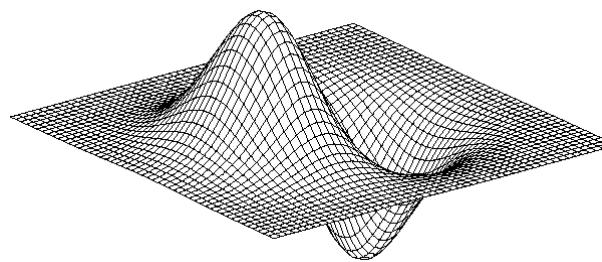
Where the function intersects with x (y=0) (**Zero-crossing**)

Summary: 2D filters for edges



Gaussian

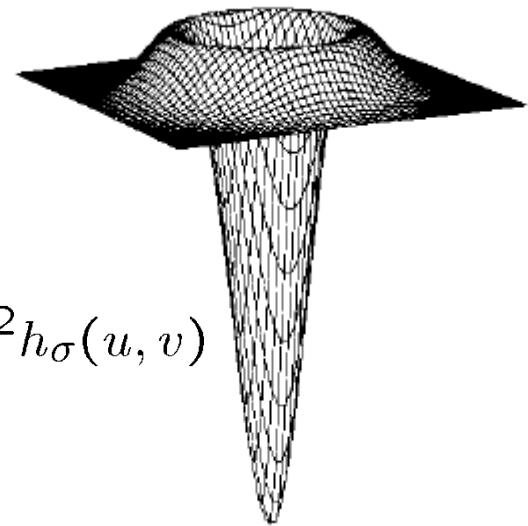
$$h_\sigma(u, v) = \frac{1}{2\pi\sigma^2} e^{-\frac{u^2+v^2}{2\sigma^2}}$$



Gaussian derivative

$$\frac{\partial}{\partial x} h_\sigma(u, v)$$

Laplacian of Gaussian

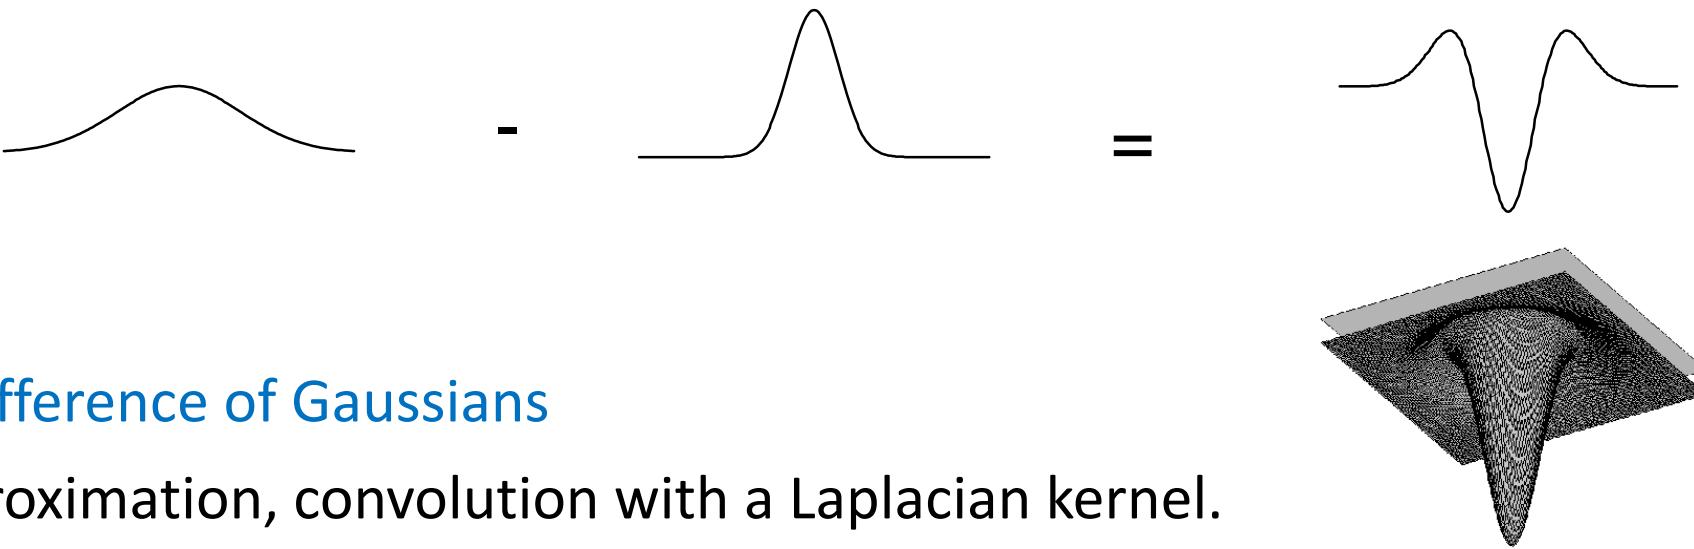


$$\nabla^2 h_\sigma(u, v)$$

- ∇^2 is a Laplace operator (a Laplacian):

$$\nabla^2 h = \frac{\partial^2 h}{\partial u^2} + \frac{\partial^2 h}{\partial v^2}$$

Laplacian \approx Difference of Gaussians



DoG = Difference of Gaussians

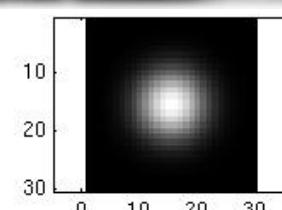
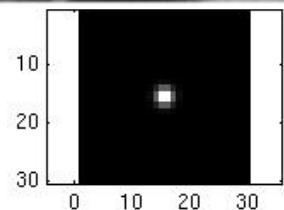
Fast approximation, convolution with a Laplacian kernel.



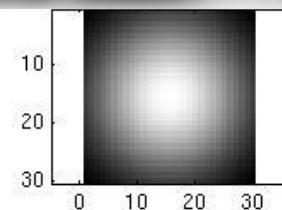
Edge scale influenced by filters

Recall:

Parameter σ is the “scale”/“width” of a Gaussian kernel that determines the extent of smoothing.

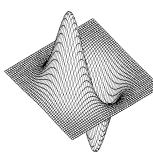


...

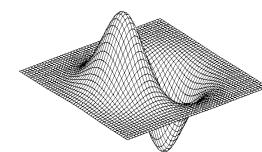
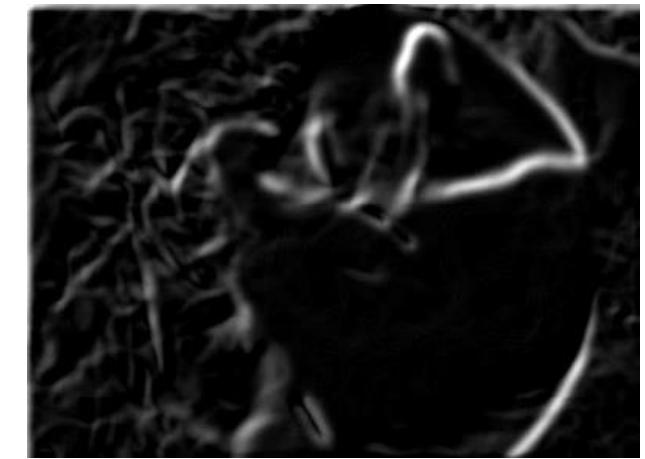


Edge scale influenced by filters

How does σ affect the derivative?



$\sigma = 1$ pixel



$\sigma = 3$ pixels

The enhanced/detected **structures** depend on the Gaussian kernel **size**.

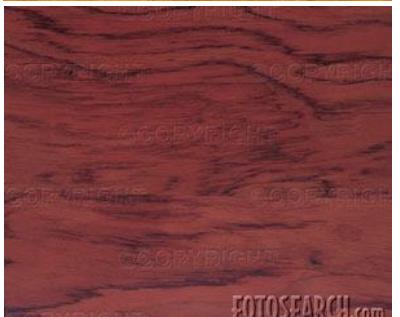
Large kernels: detect edges on a larger scale.

Small kernels: detect edges on a smaller scale.

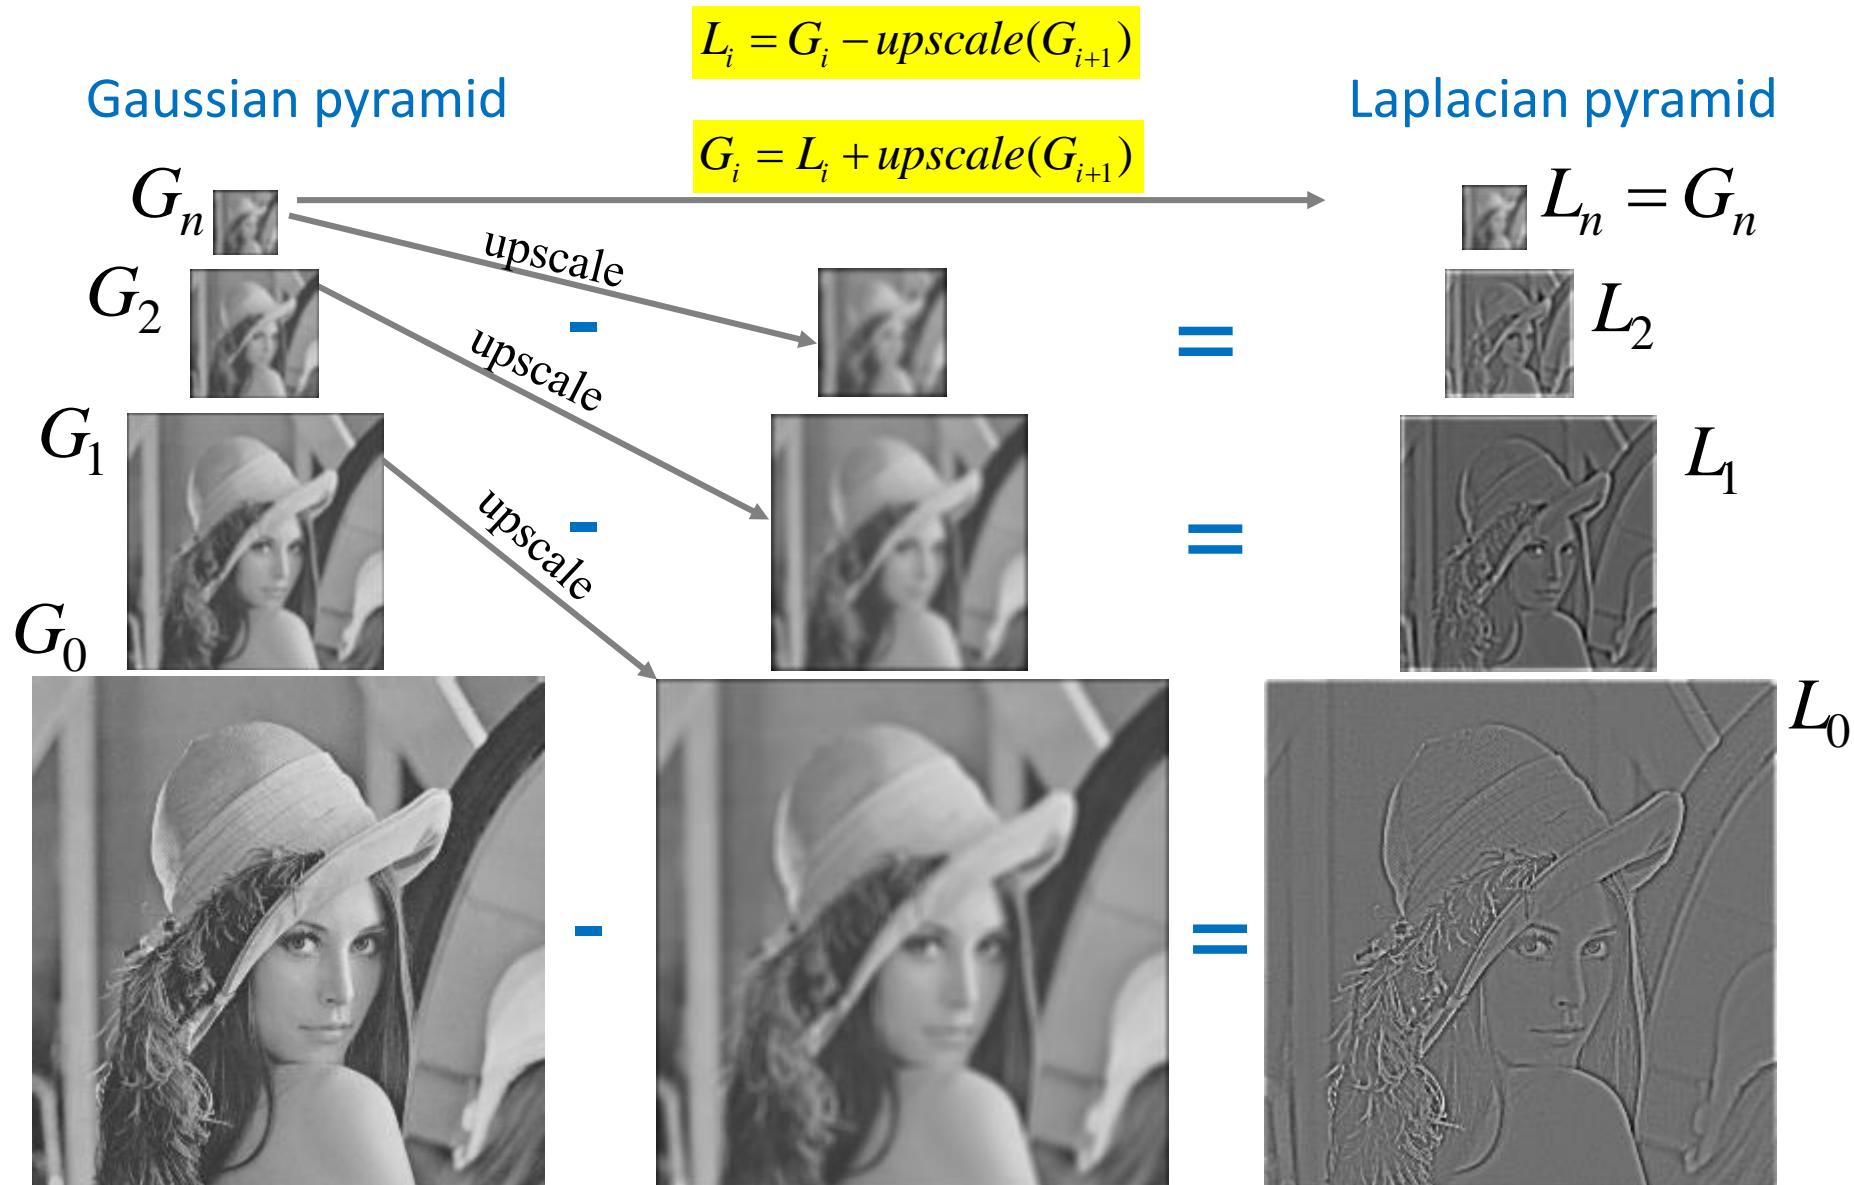
What do we mean by “detect edges”?

Depends on what we’re looking for...

Different structures are found on different scales!



Edge scale and the Laplacian pyramid



Edge scale and the Laplacian pyramid

 L_2  L_1 L_0 

- Instead of increasing the kernel size of derivative filter, the image is reduced and “filtered” with a fixed-size filter!
- Thus higher levels “respond” to thicker edges.

Machine perception

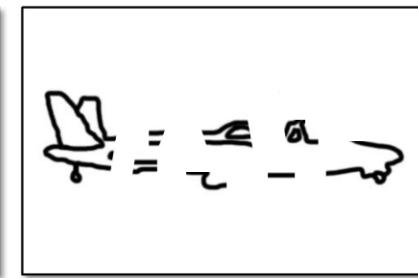
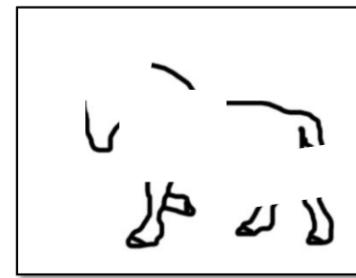
FROM DERIVATIVES TO EDGE DETECTION

The task of edge detection

- Goal: map **image** from 2D grayscale intensity pixel array **into** a set of binary **curves** and **lines**.
- Why?



abstraction



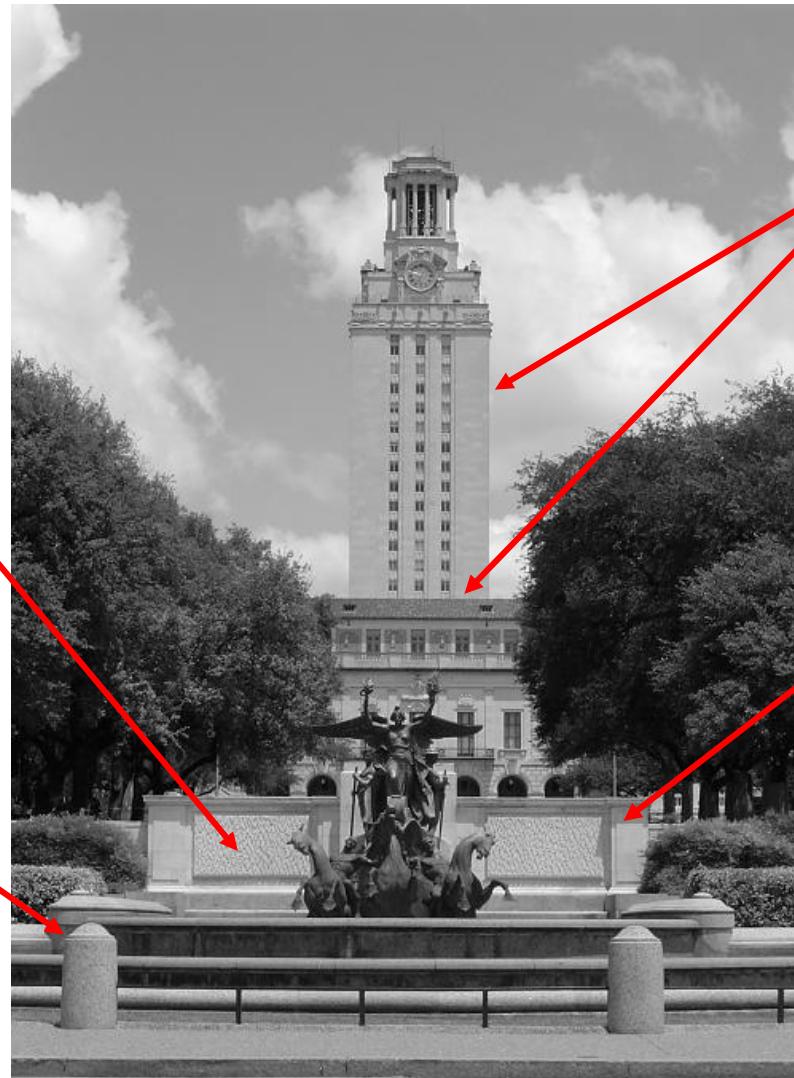
Robust, compact representation

What constitutes an edge?

Anything that *appears*
as an edge...

Local texture:

Changes in 3D normal
orientation caused by
shape changes

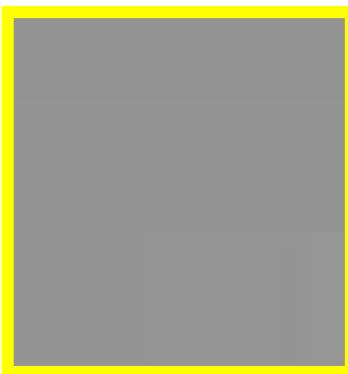


Discontinuity of depth:
object borders

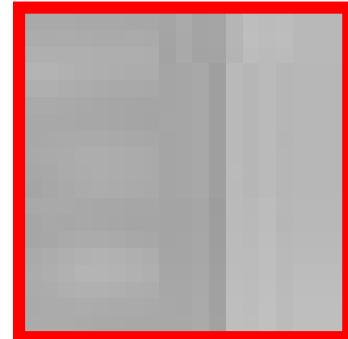
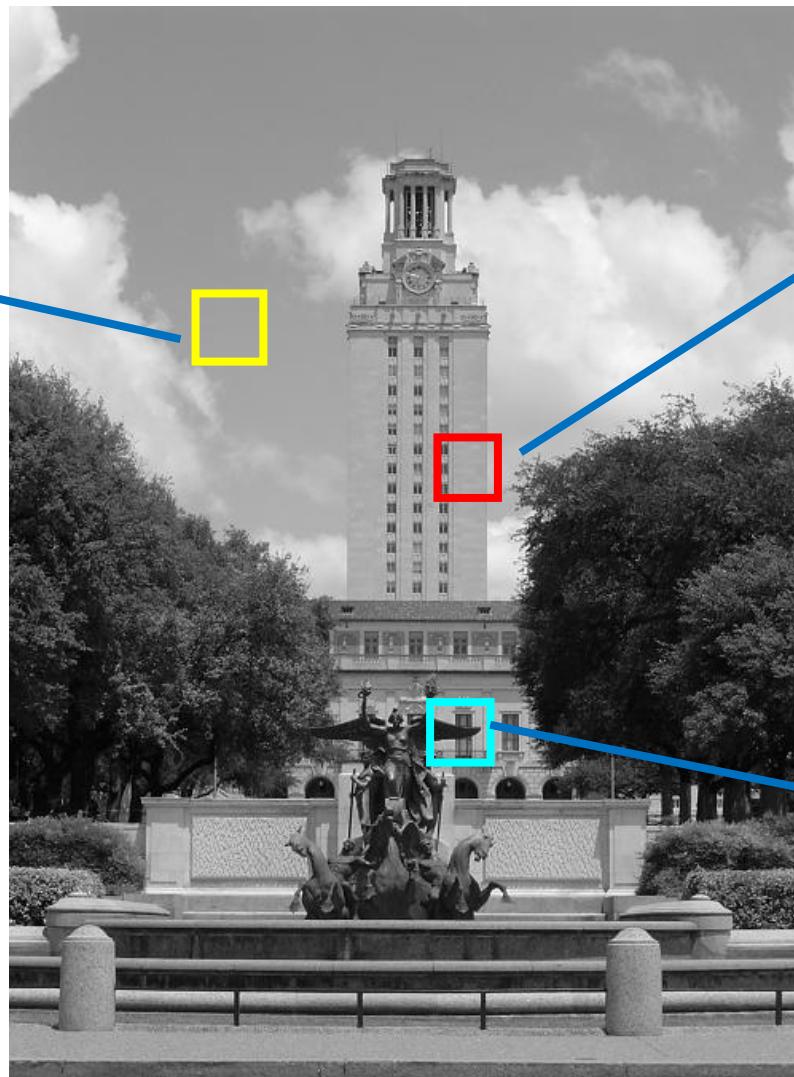
Shadows

What constitutes an edge?

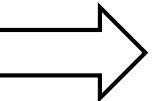
Anything that *appears*
as an edge...



Edge presence is strongly
correlated with the local
intensity changes.

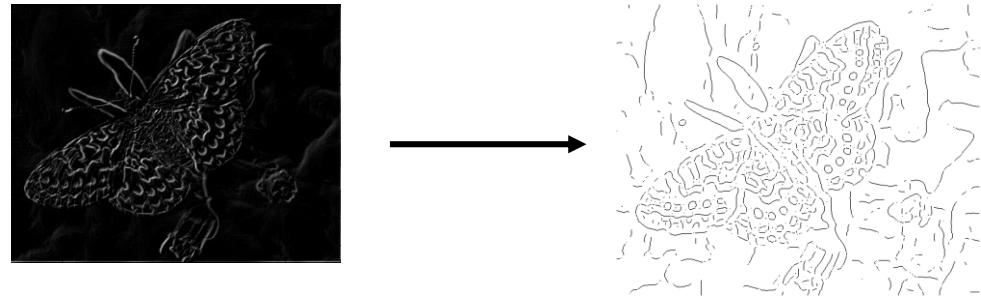


The task of edge detection



- Basic approach:
find strong **gradients** + post process

From gradients to edges



Three steps for edge detection

1. Smoothing:

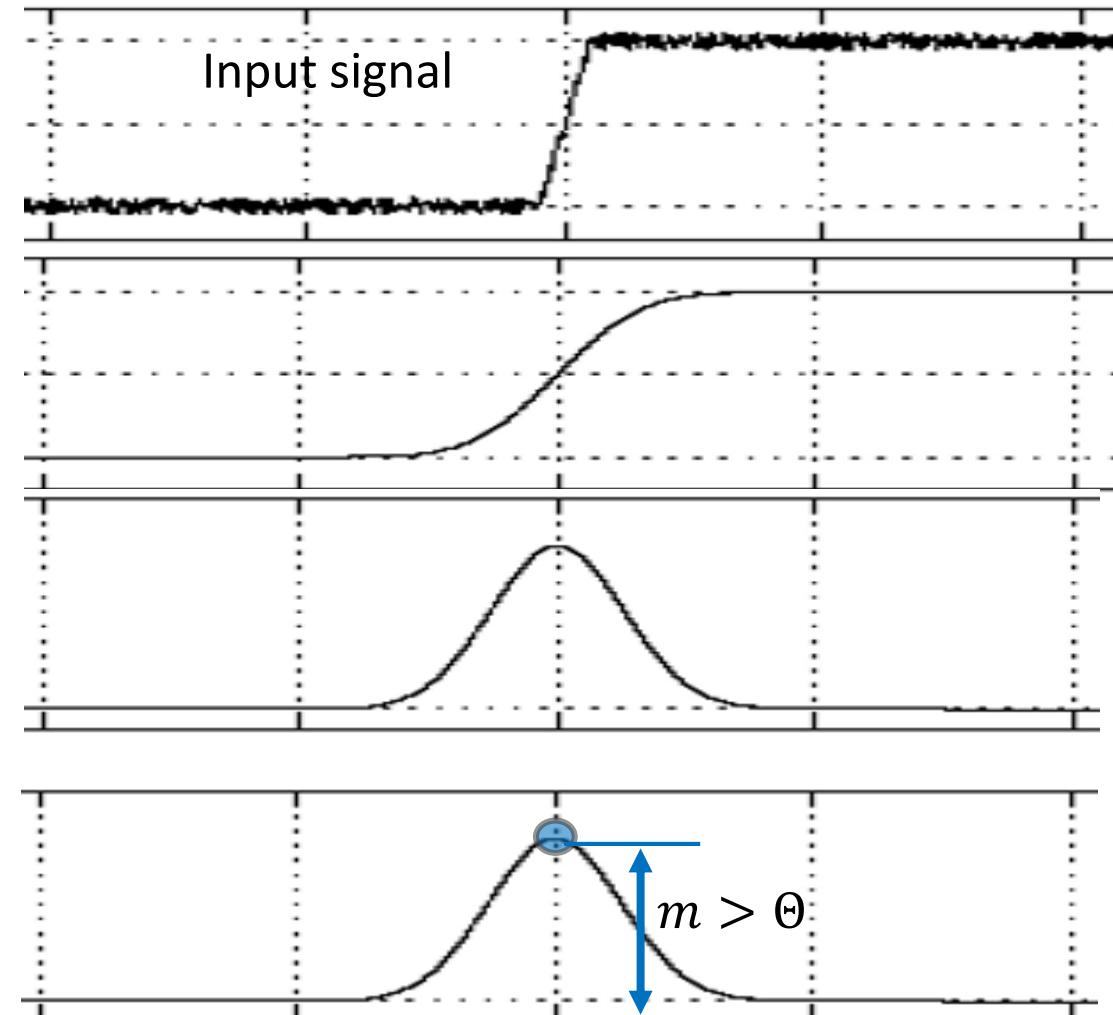
noise suppression

2. Edge enhancing:

increase contrast (by differentiating)

3. Edge localization

- Identify potential edge pixels
- Verify the magnitude

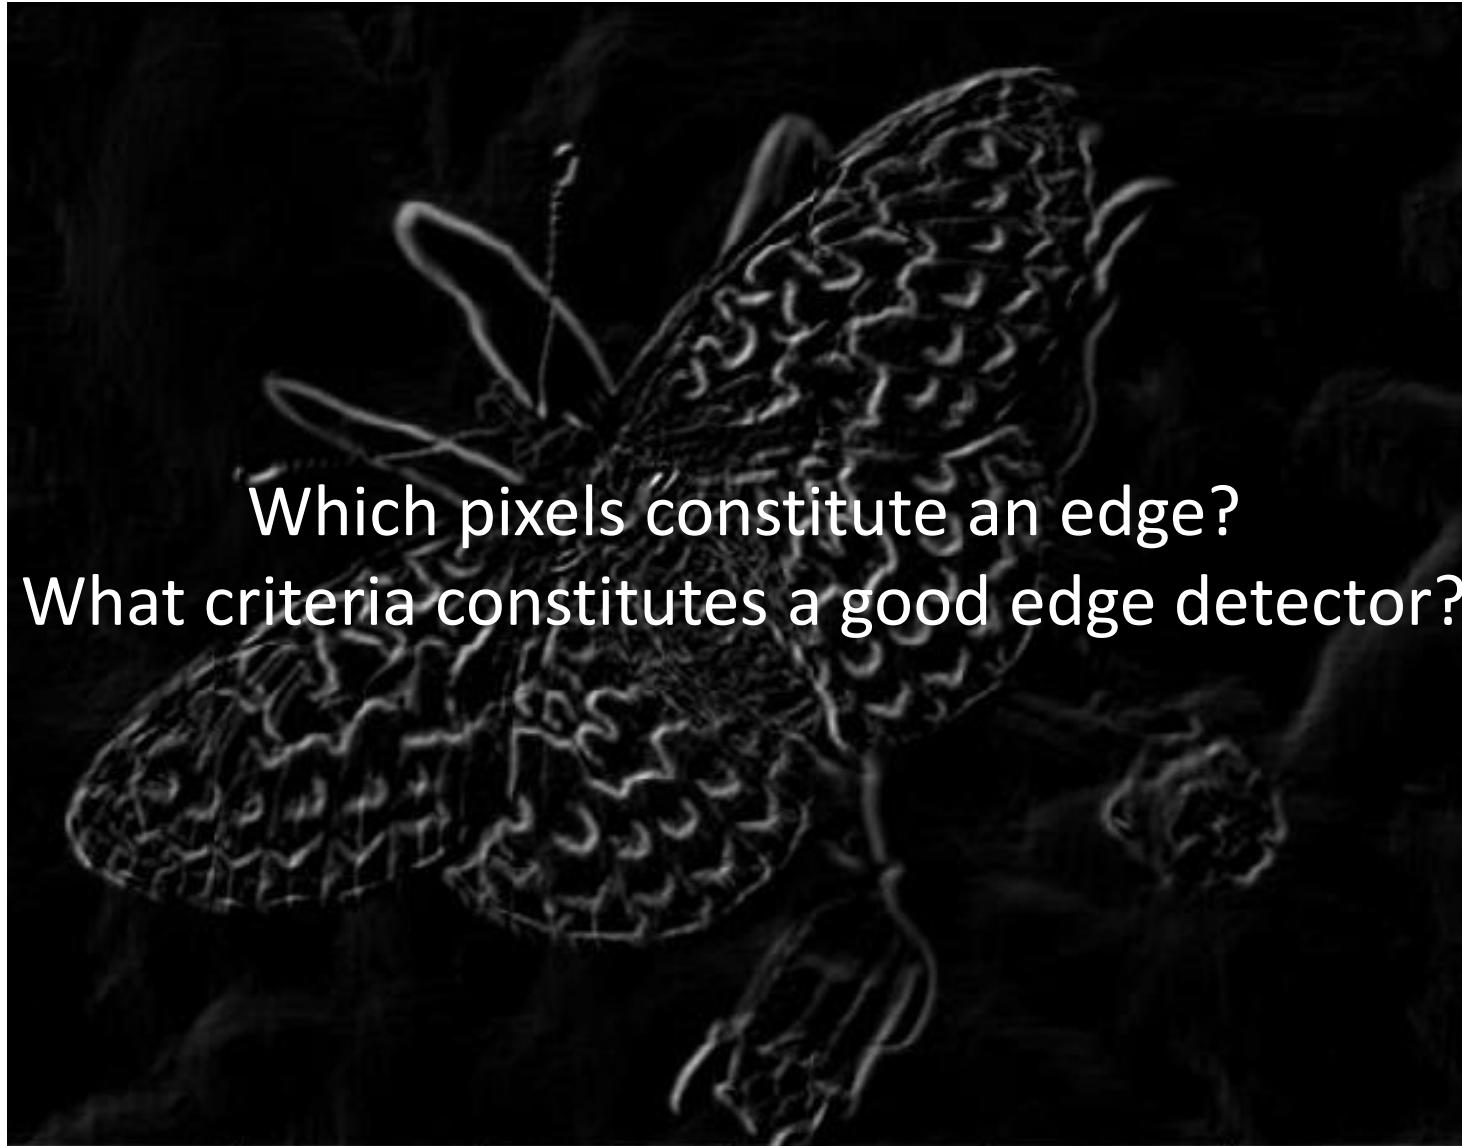


Original image



Slide credit: Kristen Grauman

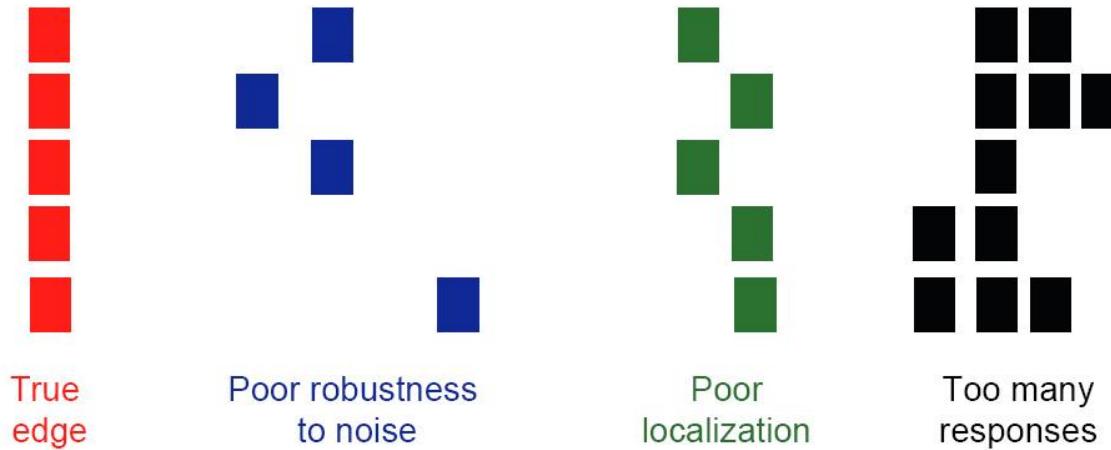
Gradient magnitudes (smooth+enhance)



Which pixels constitute an edge?
What criteria constitutes a good edge detector?

Designing an edge detector...

- Criteria of “optimal” edge detector:
 1. **Good detection:** optimal detector **minimizes** probability of **false positives** (edges caused by noise), and **false negatives** (missing true edges)
 2. **Good localization:** detected edges should be **close** to the location of the **true edges**.
 3. **Specificity:** detector should return only **a single point** per **true edge**; minimize number of local maxima around true edge.



The Canny edge detector

- Most popular edge detector in computer vision.
- Theoretical model:
Step edges with added Gaussian noise.
- Canny showed that first derivative of a Gaussian well approximates an operator that optimizes a tradeoff between signal-to-noise ratio and localization.

MATLAB:

```
>> edge(image, 'canny');  
>> help edge
```

Canny edge detector

1. Filter image by a derivative of a Gaussian
2. Calculate the gradient magnitude and orientation
3. Non-maxima suppression:
 - Convert multiple pixels-thick edges to a single-pixel thickness: Thinning
 -

$$\nabla f = \left[\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right]$$

$$\theta = \tan^{-1} \left(\frac{\partial f}{\partial y} / \frac{\partial f}{\partial x} \right)$$

$$\|\nabla f\| = \sqrt{\left(\frac{\partial f}{\partial x}\right)^2 + \left(\frac{\partial f}{\partial y}\right)^2}$$

Canny edge detector: Thinning



Original image (Lena)

Canny edge detector: Thinning – 1. Derivative



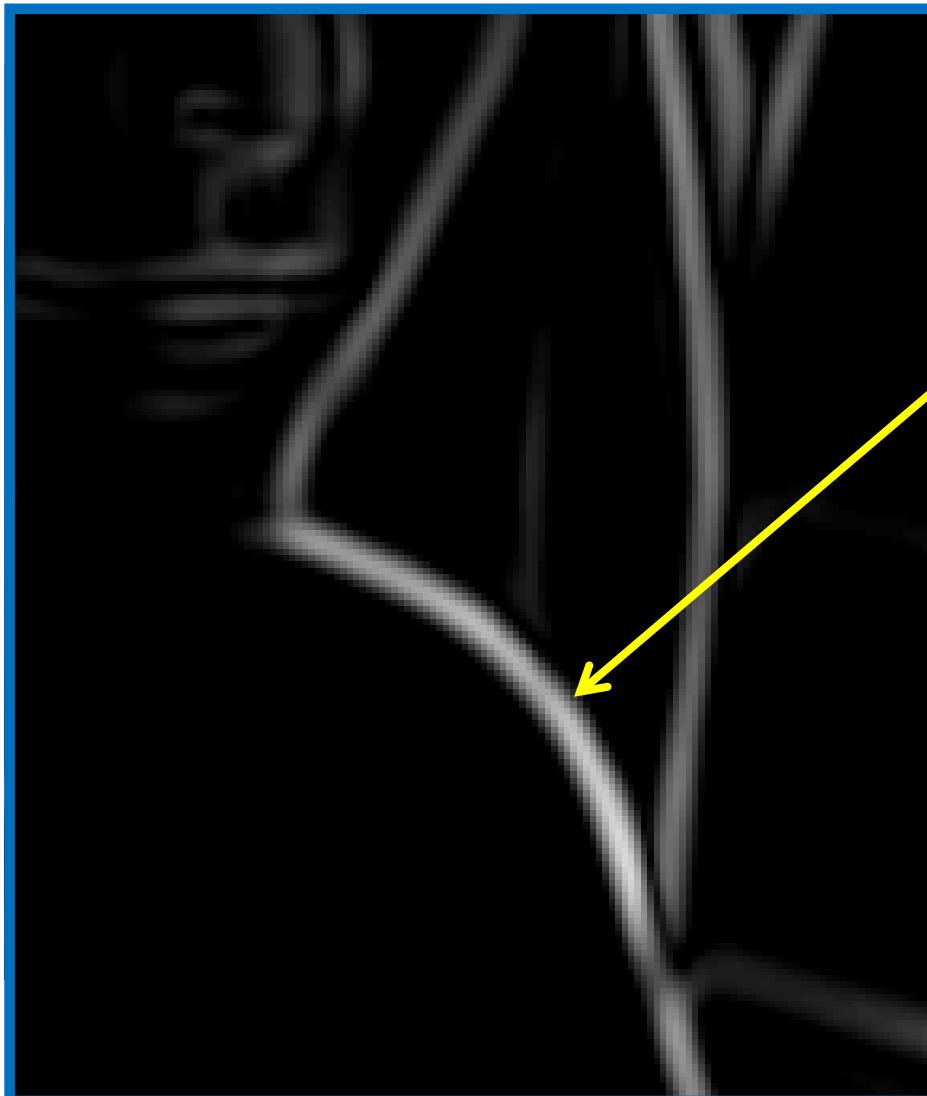
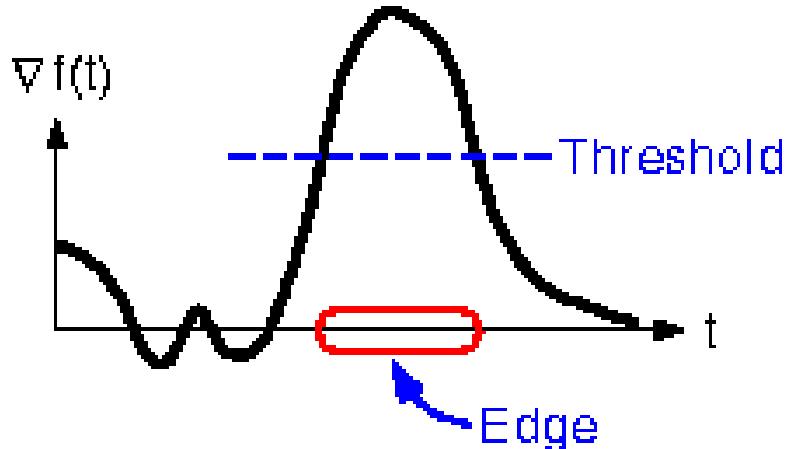
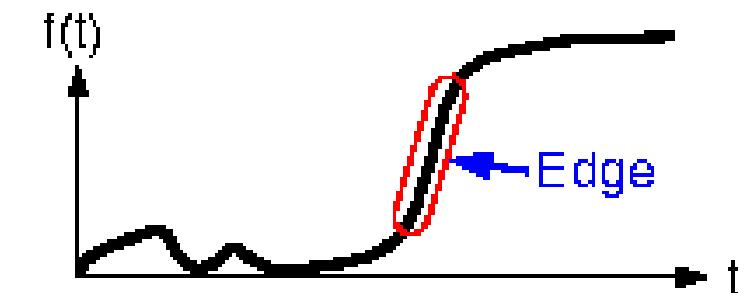
Gradient magnitude

Canny edge detector: Thinning - 2. Threshold



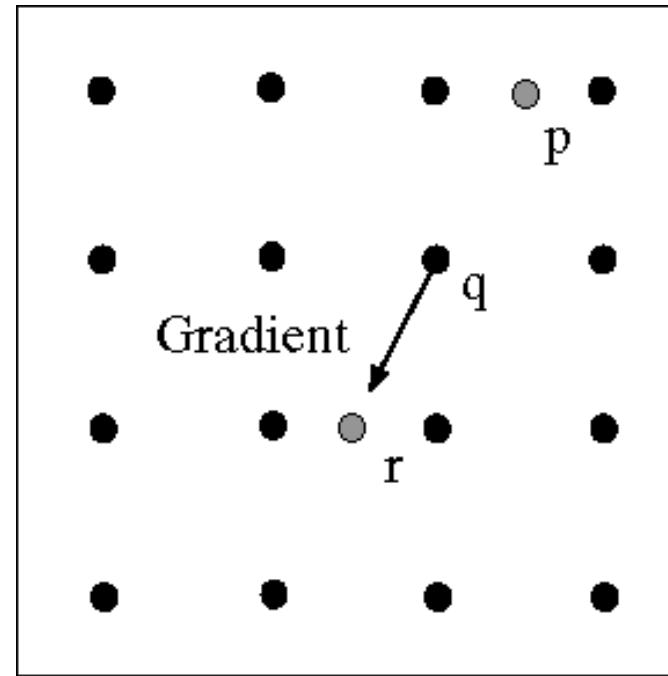
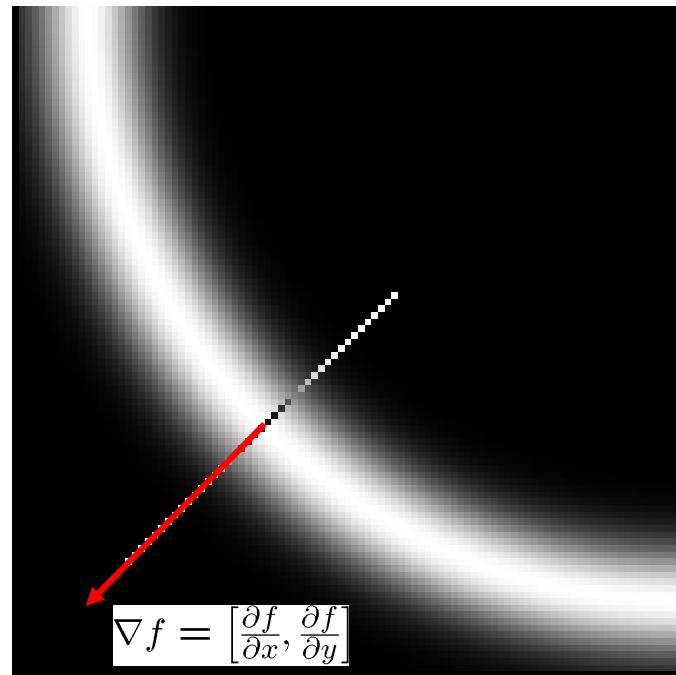
„Thresholding“: Set magnitudes lower than a prescribed threshold to 0.

Canny edge detector: Thinning



How to convert
these thick
lines into
thinner curves?

Thinning by non-maxima suppression



- For each pixel check if it is a **local maximum** along its gradient direction.
- Only local maxima should remain.
 - Advanced: Actually, for **q**, we should check **interpolated** pixel values at **p** and **r**.

Canny edge detector: Thinning – 3. Nonmax Suppression



Thinning
(non-maximum suppression)

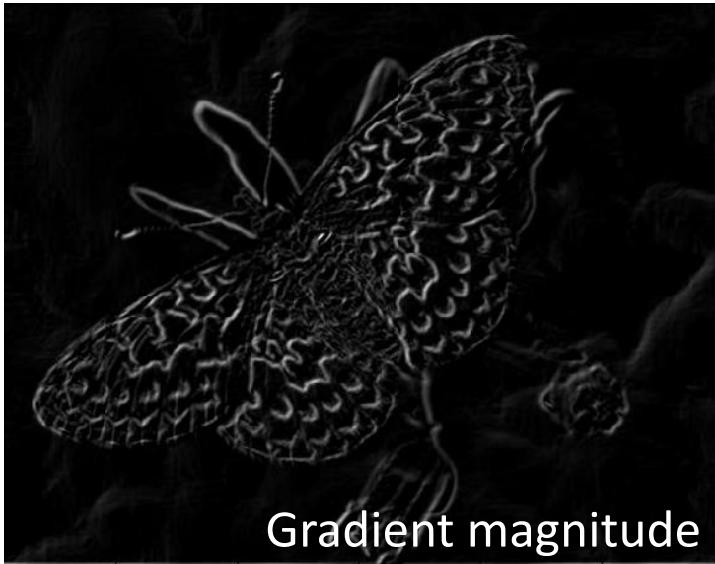
Canny edge detector: Thinning



Thinning
(non-maximum suppression)

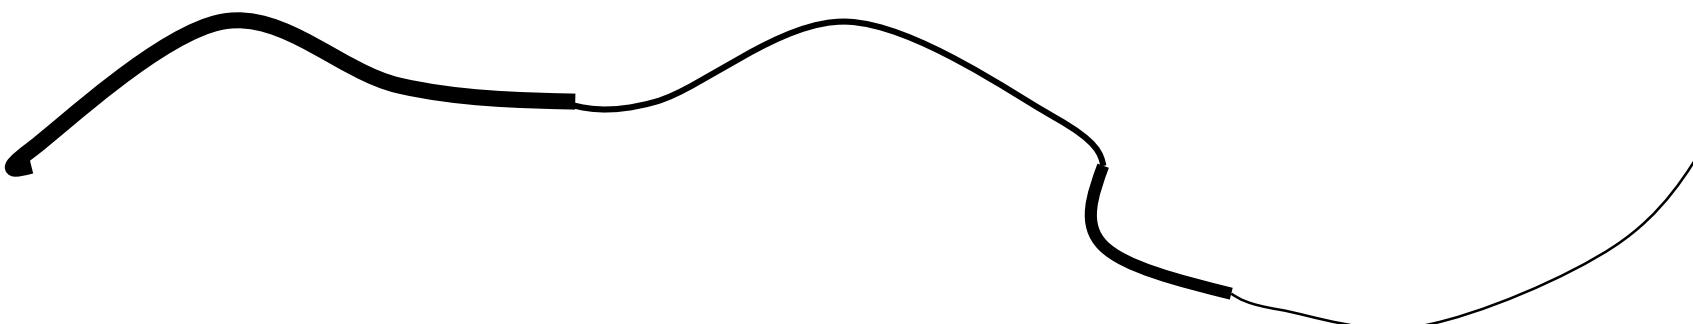
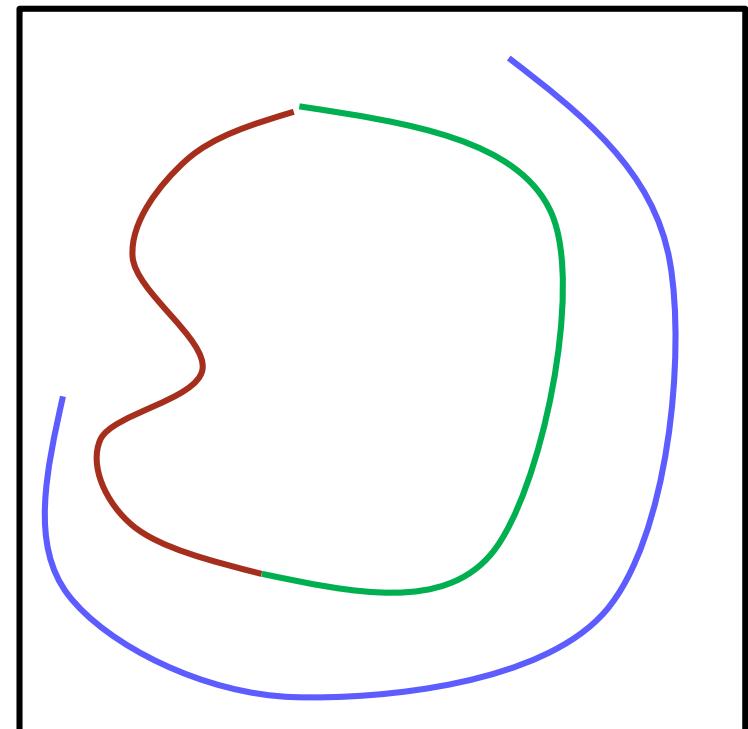
Problem: pixels
along this edge
did not „survive“
the thresholding.

How to select a threshold?



Canny edge detector: Hysteresis thresholding

- Trace each contour separately (e.g., using 4-connectedness).
- Apply two thresholds k_{high} and k_{low}
 - Start tracing a line only at pixels that pass the k_{high} threshold.
 - Continue tracing if the pixels pass k_{low} threshold.
- Typical threshold ratio: $k_{\text{high}} / k_{\text{low}} = 2$



Hysteresis thresholding



Original



High threshold
(strong edges)



Low threshold
(weak edges)



Hysteresis thresholding

The Canny edge detector in a nutshell

1. Convolve the image by a derivative of a Gaussian.
2. Calculate the gradient magnitude and orientation
3. Non-maxima suppression
 - Thin edges to one-pixel width.
4. Trace the edges by hysteresis thresholding
 - Apply a high threshold on the magnitude to initialize a contours and continue tracing the contour until the magnitude falls below a low threshold.

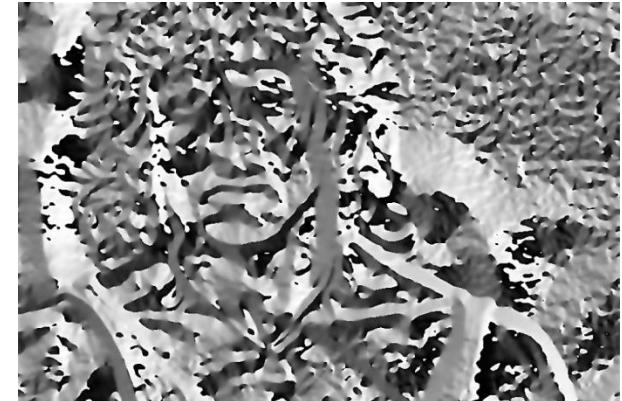
Canny edge detector in “action”



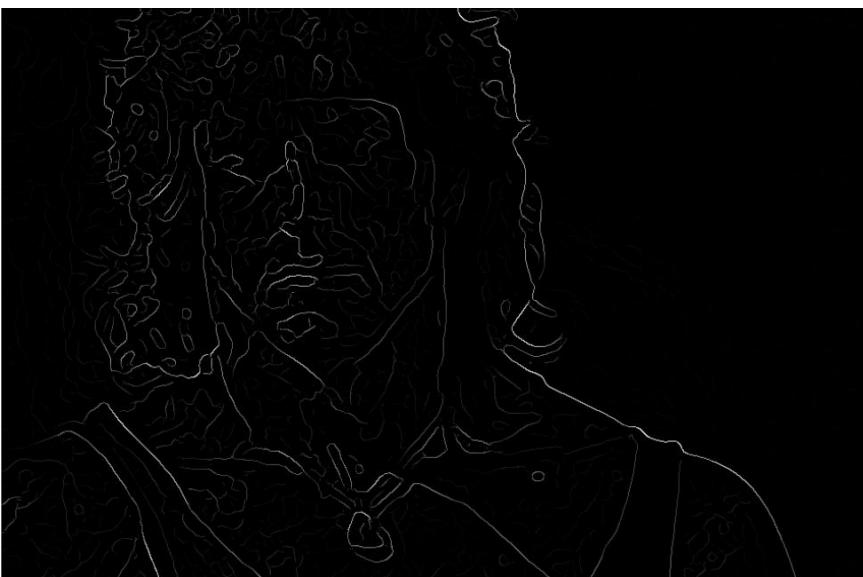
Input image



gradient magnitude



gradient angle



thinned



Thresholded by hysteresis

Canny edge detector in “action”



Input image



gradient magnitude



gradient angle



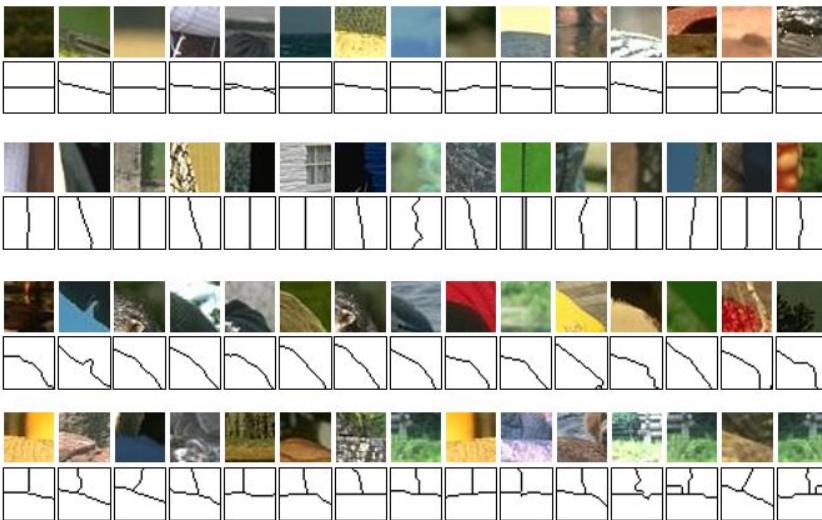
thinned



Thresholded by hysteresis

Beyond Canny edge detector

- Recently lots of new approaches for edge detection by machine learning.
- Essentially, look at patches and learn what an edge is by inferring the structure from intensities.



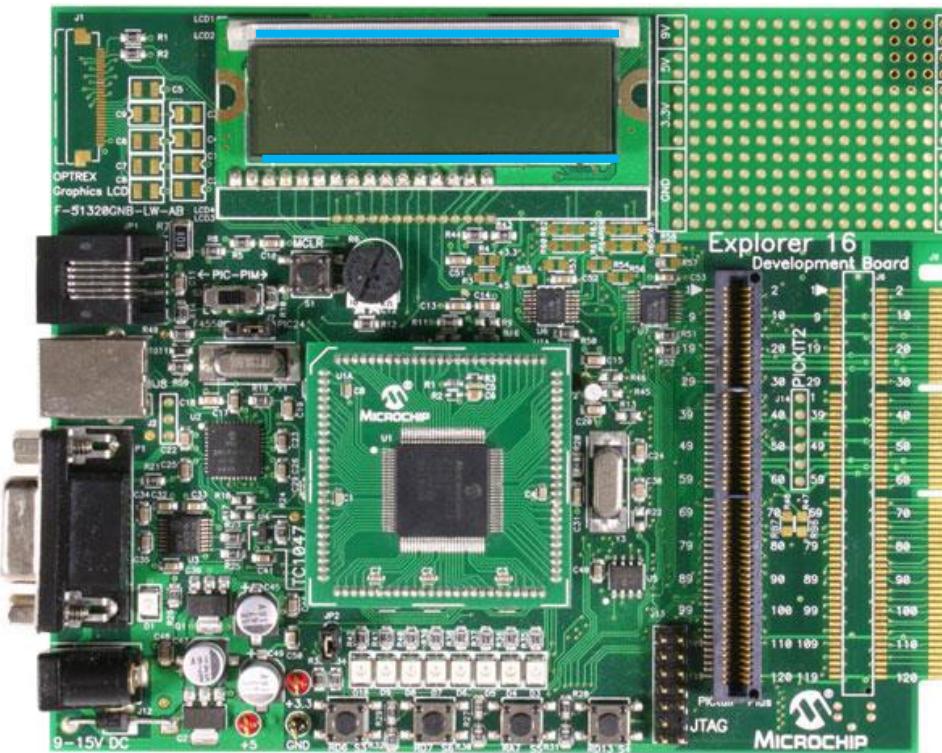
Machine perception

EDGE DETECTION USING PARAMETRIC MODELS

Example: line fitting

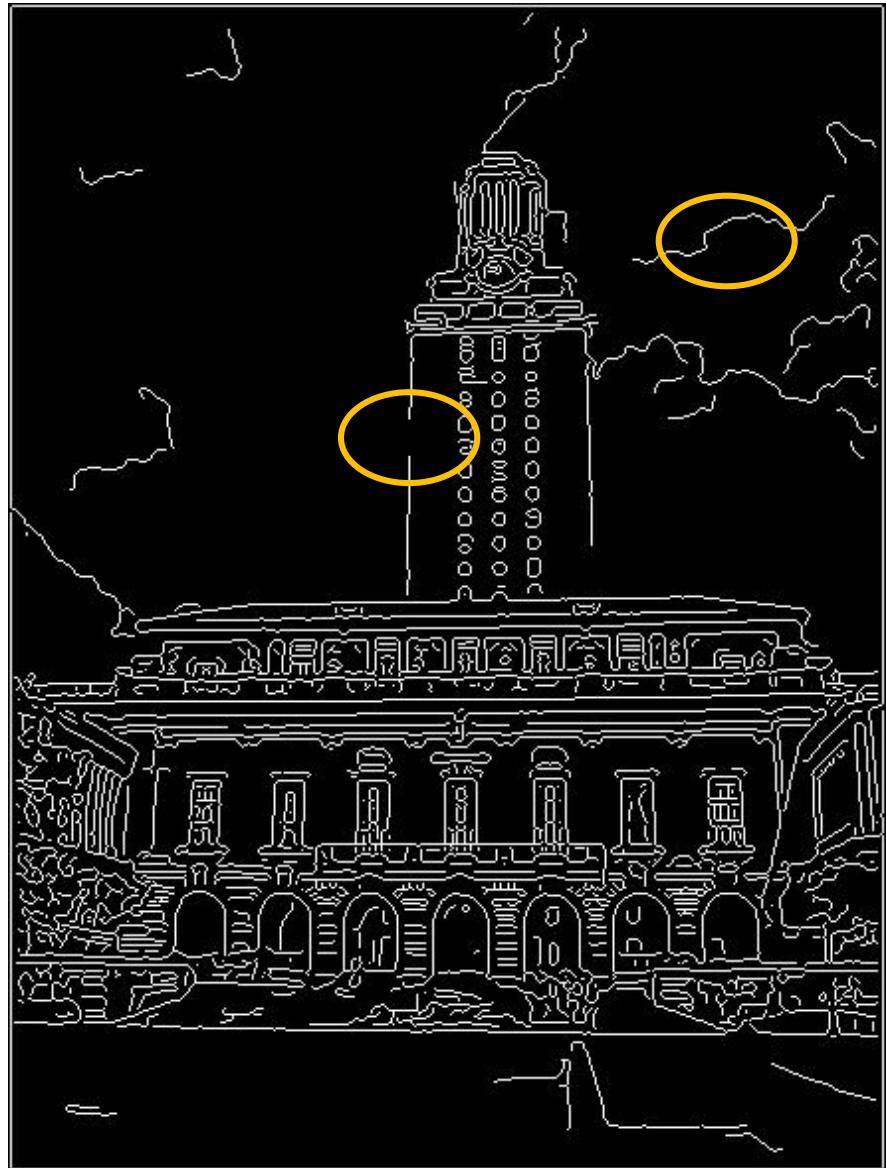
- Why should we fit lines?

Many scenes are composed of straight lines

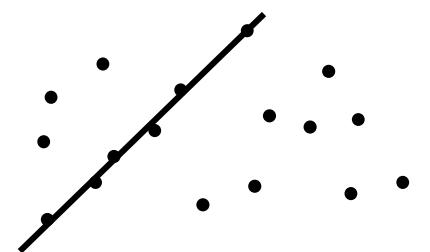


- But, haven't we done that already by edge detection?

Problems with line fitting

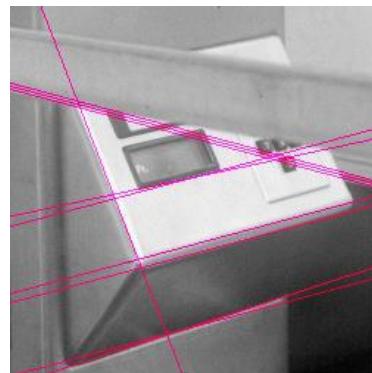
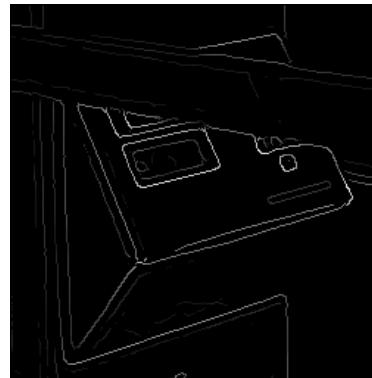


- Noisy edges, multiple models:
 - Which points correspond with which line, if at all?
- Some parts of lines are not detected:
 - How to find a line that connects the missing points?
- Noisy orientation:
 - How do we determine the unknown parameters of true lines?

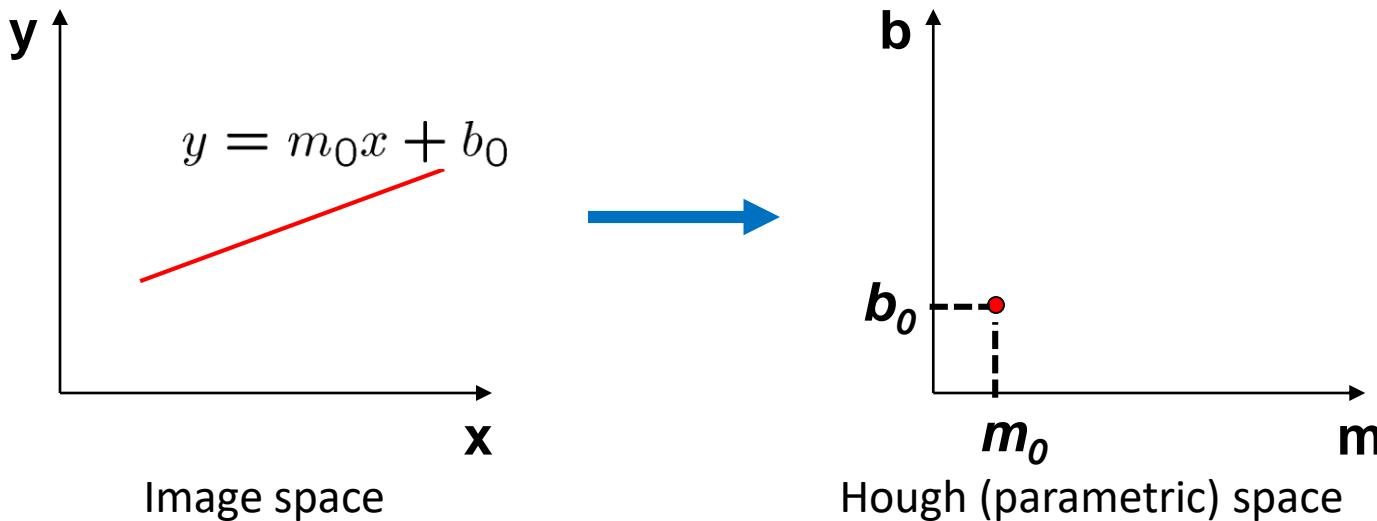


Line fitting by voting for parameters

- Given a set of points, find the lines.
 - How many lines?
 - Which points correspond to which lines?
-
- *Hough Transform* is a voting technique that answers these questions.
 - Main idea:
 1. For each edge point compute all possible lines that pass through that point.
 - Each point casts a vote for parameter values.
 2. Select the lines (parameter combinations) that receive enough votes.

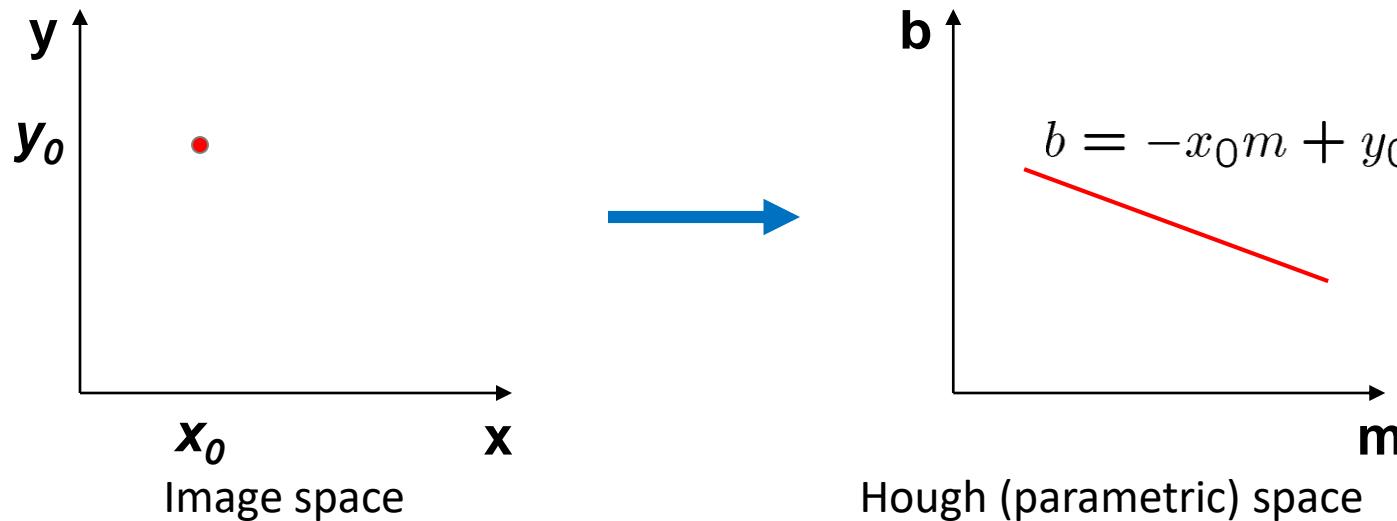


Hough space: straight lines



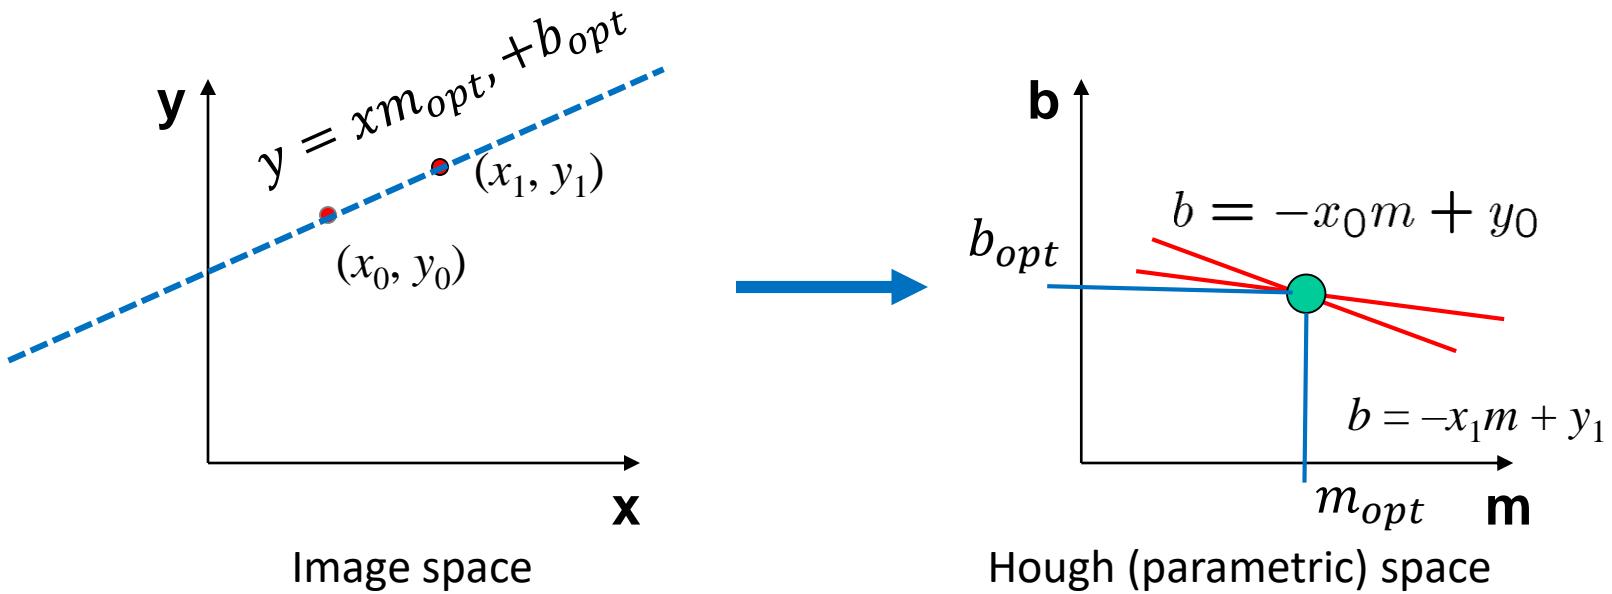
- Connection between spatial (x,y) and Hough space (m,b) :
 - A line in image corresponds to a point in the Hough space.

Hough space: straight lines



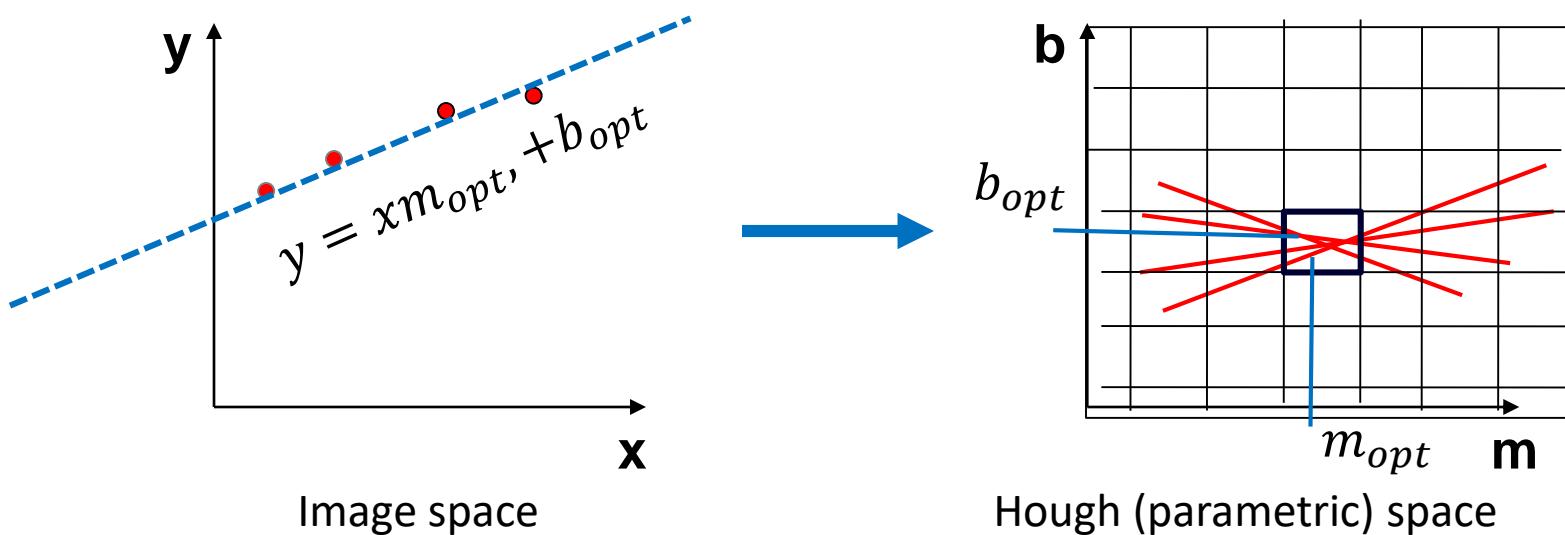
- Connection between spatial (x,y) and Hough space (m,b) :
 - A line in image corresponds to a point in the Hough space.
 - Mapping from image to Hough space:
 - For a point (x,y) , find all (m,b) for which this holds : $y = mx + b$

Hough space: straight lines



- Connection between spatial (x,y) and Hough space (m,b) :
 - A line in image corresponds to a point in the Hough space.
 - Mapping from image to Hough space:
 - For a point (x,y) , find all (m,b) for which this holds : $y = mx + b$

Hough space: straight lines

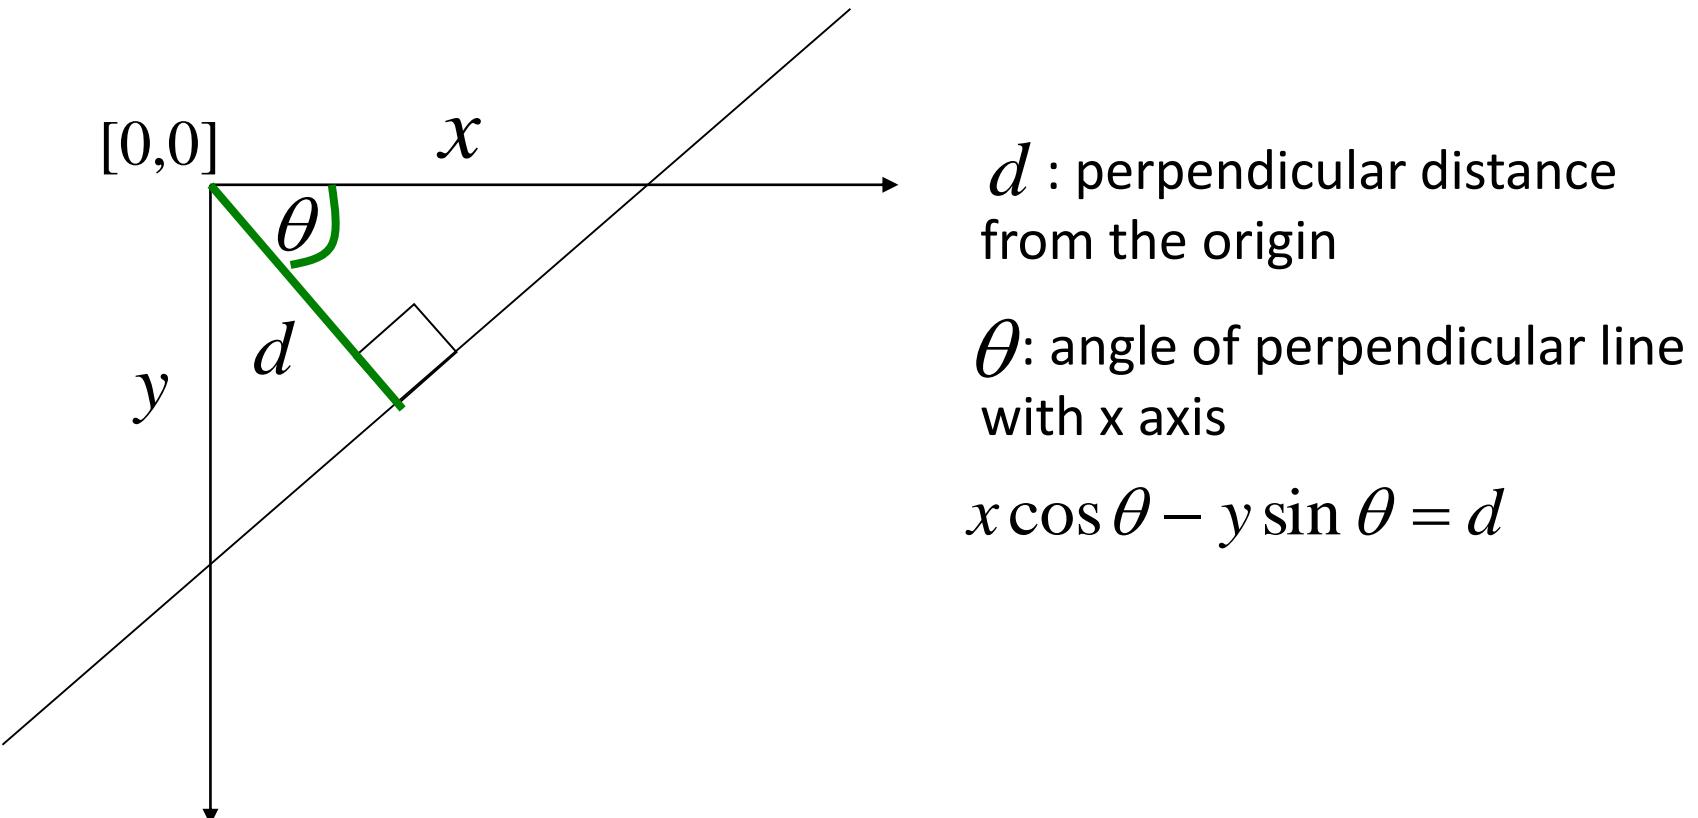


Discretize the parameter space...

- Connection between spatial (x,y) and Hough space (m,b) :
 - A line in image corresponds to a point in the Hough space.
 - Mapping from image to Hough space:
 - For a point (x,y) , find all (m,b) for which this holds : $y = mx + b$

Encode the line in polar coordinates

- Issue with Cartesian (m,b) : infinite values for vertical lines!



- Point in image \Rightarrow sinusoid in Hough space

Algorithm: Straight lines

Using polar representation:

$$x \cos \theta - y \sin \theta = d$$

Basic Hough transform:

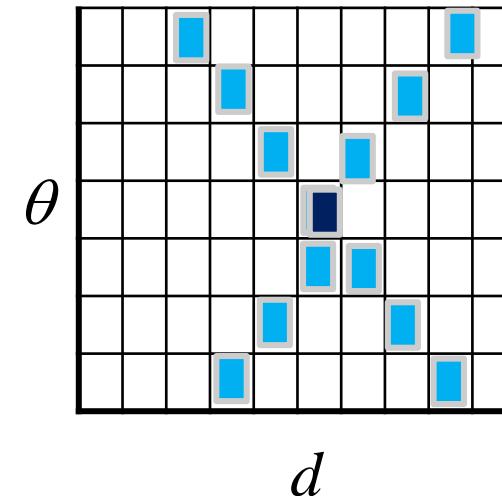
1. Initialize $H[d, \theta] = 0$.
2. For each edge point (x, y) in image
For $\theta = 0$ to 180 // over quantized values!!

$$d = x \cos \theta - y \sin \theta$$

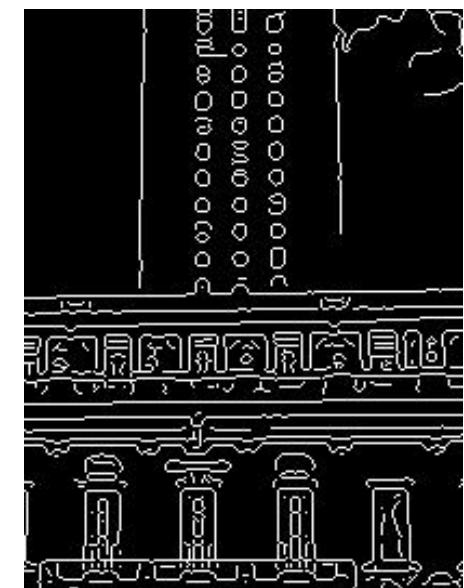
$$H[d, \theta] += 1$$

3. Find local maxima (d, θ) in accumulator array $H[d, \theta]$.
4. Detected line is defined by: $d = x \cos \theta - y \sin \theta$

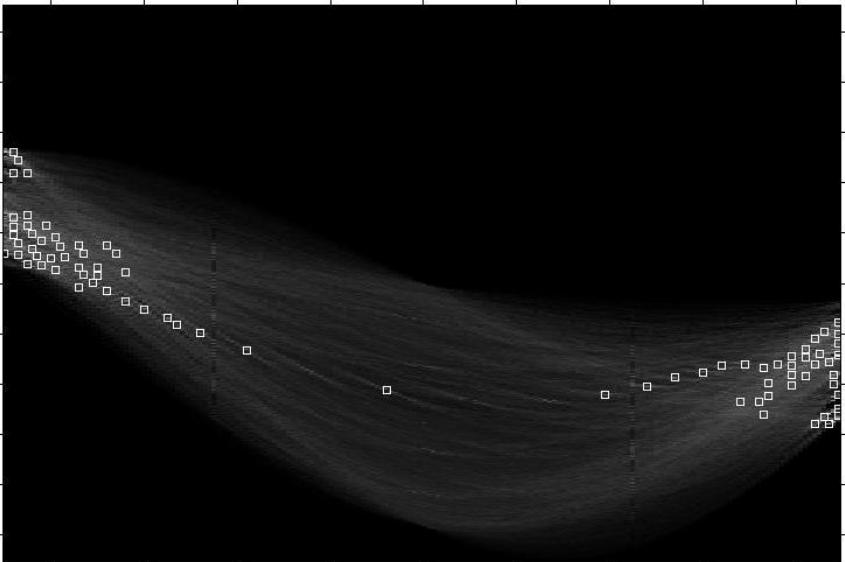
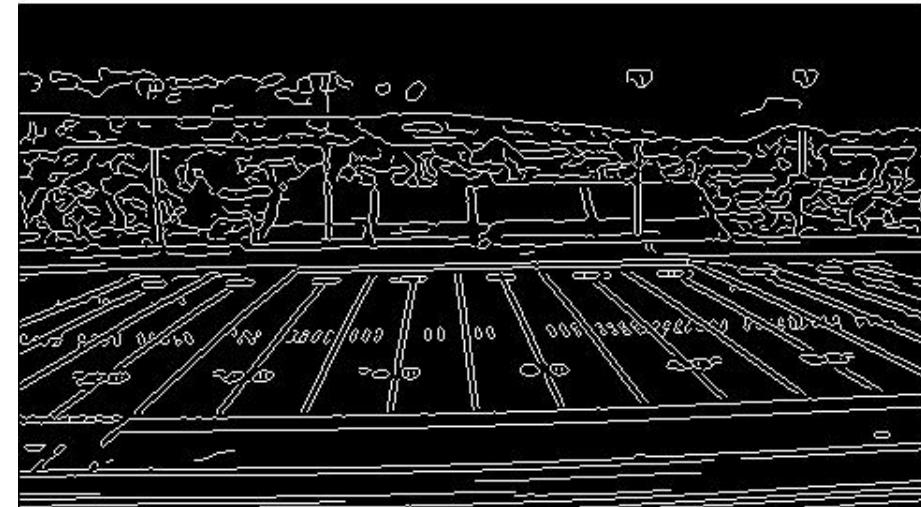
H : accumulator array (votes)



[Hough line demo](#)

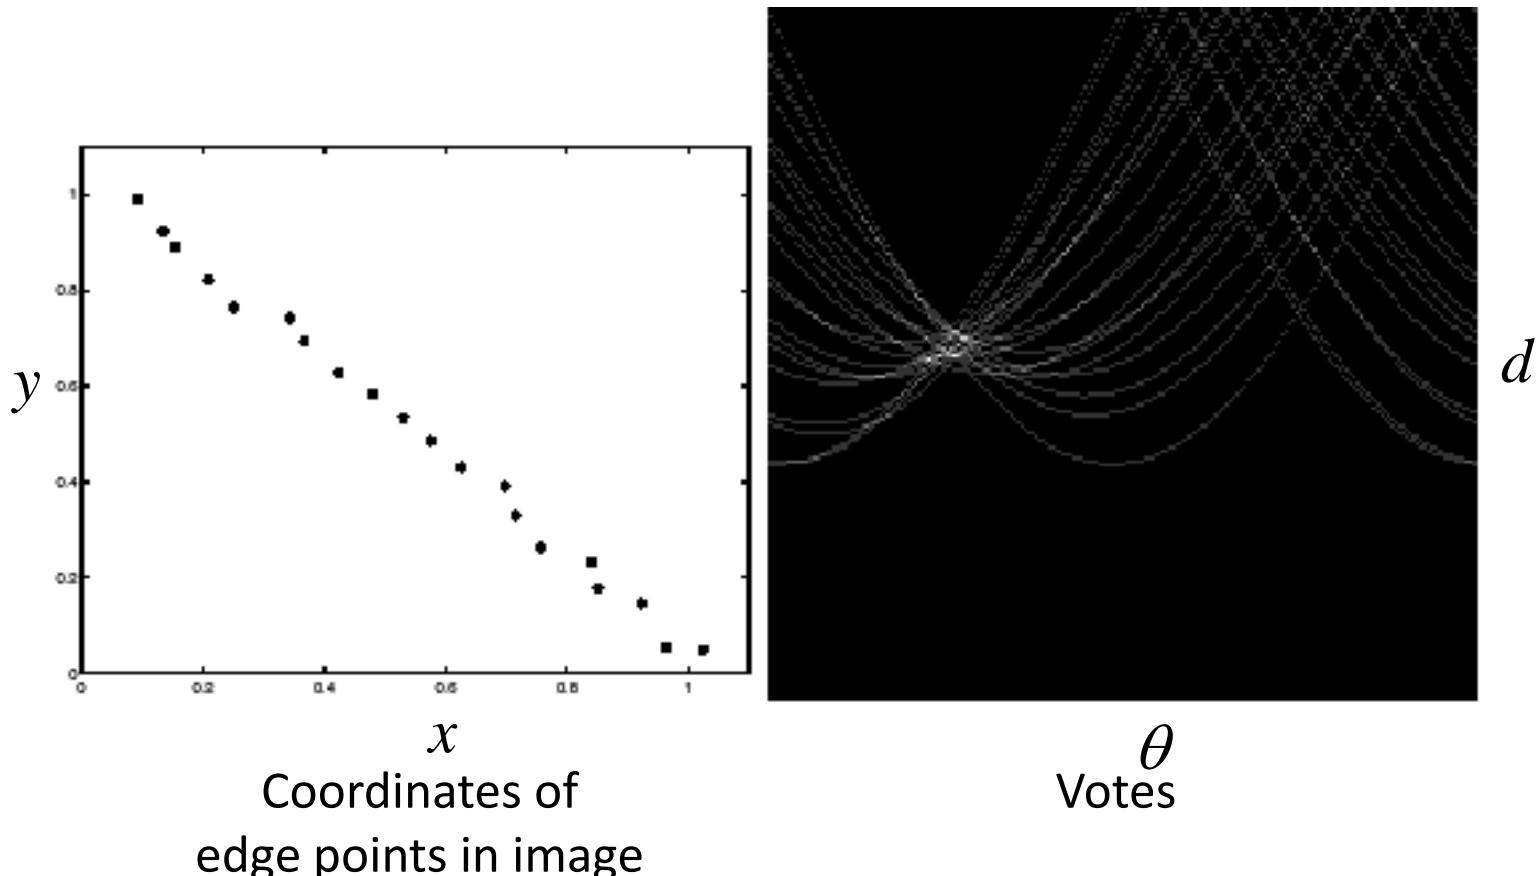


Hough transform in action



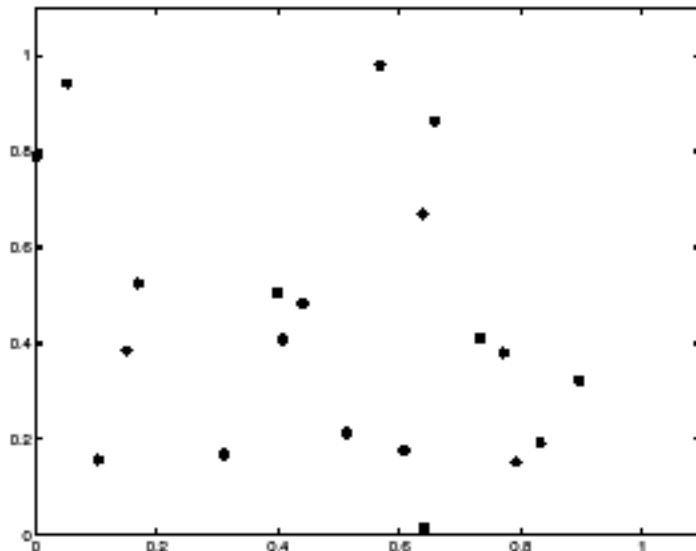
Only the longest segments shown here.

Hough transform: Noise – binning

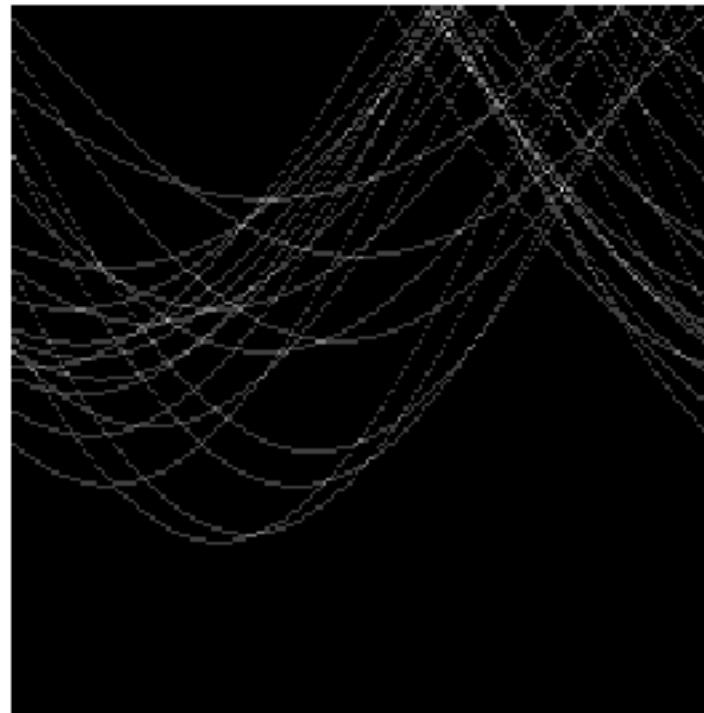


Are there any significant **problems** with the noise?

Hough transform: Noise – amplitude of votes



Coordinates of
edge points in image



Votes

Random points still **form some local maxima** in the accumulator array!

Hough transform: Extensions

Extension 1: Use the gradient info!

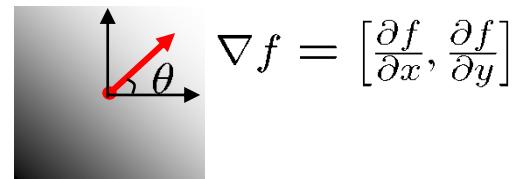
1. same as standard HT
2. For each edge point $[x, y]$
 θ = gradient at (x, y)

$$d = x \cos \theta - y \sin \theta$$

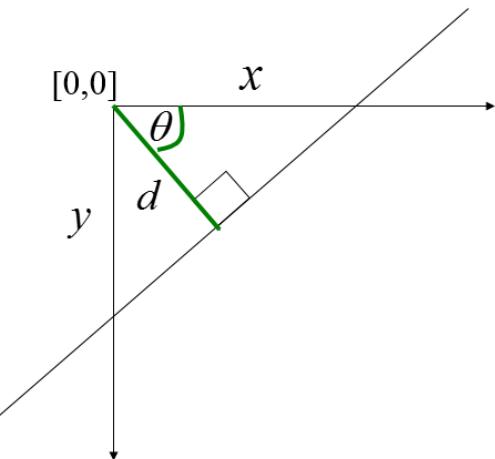
$$H[d, \theta] += 1$$

3. same as standard HT
4. same as standard HT

Reduces the number of degrees of freedom (dof)!



$$\theta = \tan^{-1} \left(\frac{\partial f / \partial y}{\partial f / \partial x} \right)$$



Hough transform: Extensions

Extension 1: Use the gradient info!

1. same as standard HT

2. For each edge point $[x, y]$

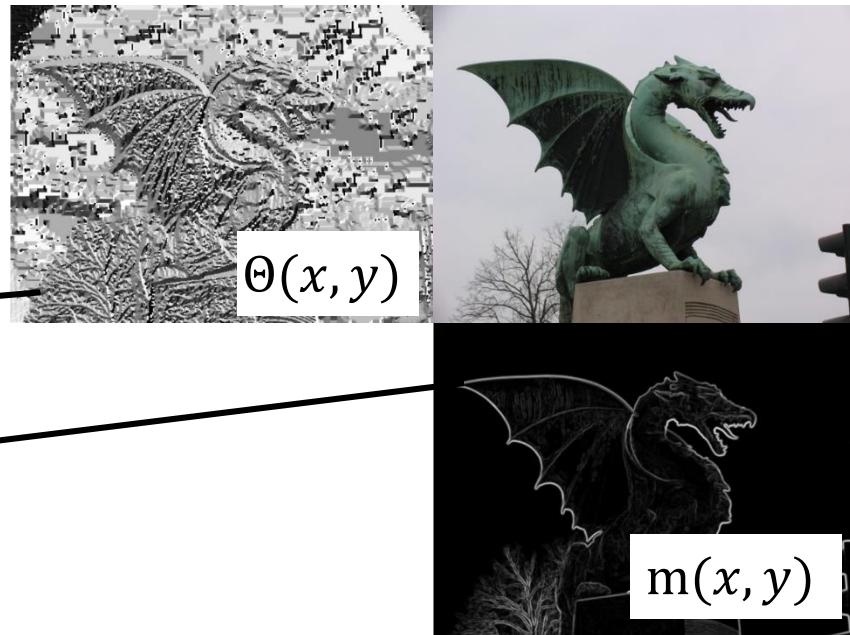
$$\theta = \text{gradient at } (x, y)$$

$$d = x \cos \theta - y \sin \theta$$

$$H[d, \theta] += 1$$

3. same as standard HT

4. same as standard HT



Extension 2:

- Assign higher weight in votes to points with large edge magnitude.

Instead $H[d, \theta] += 1$, use $H[d, \theta] += m(x, y)$.

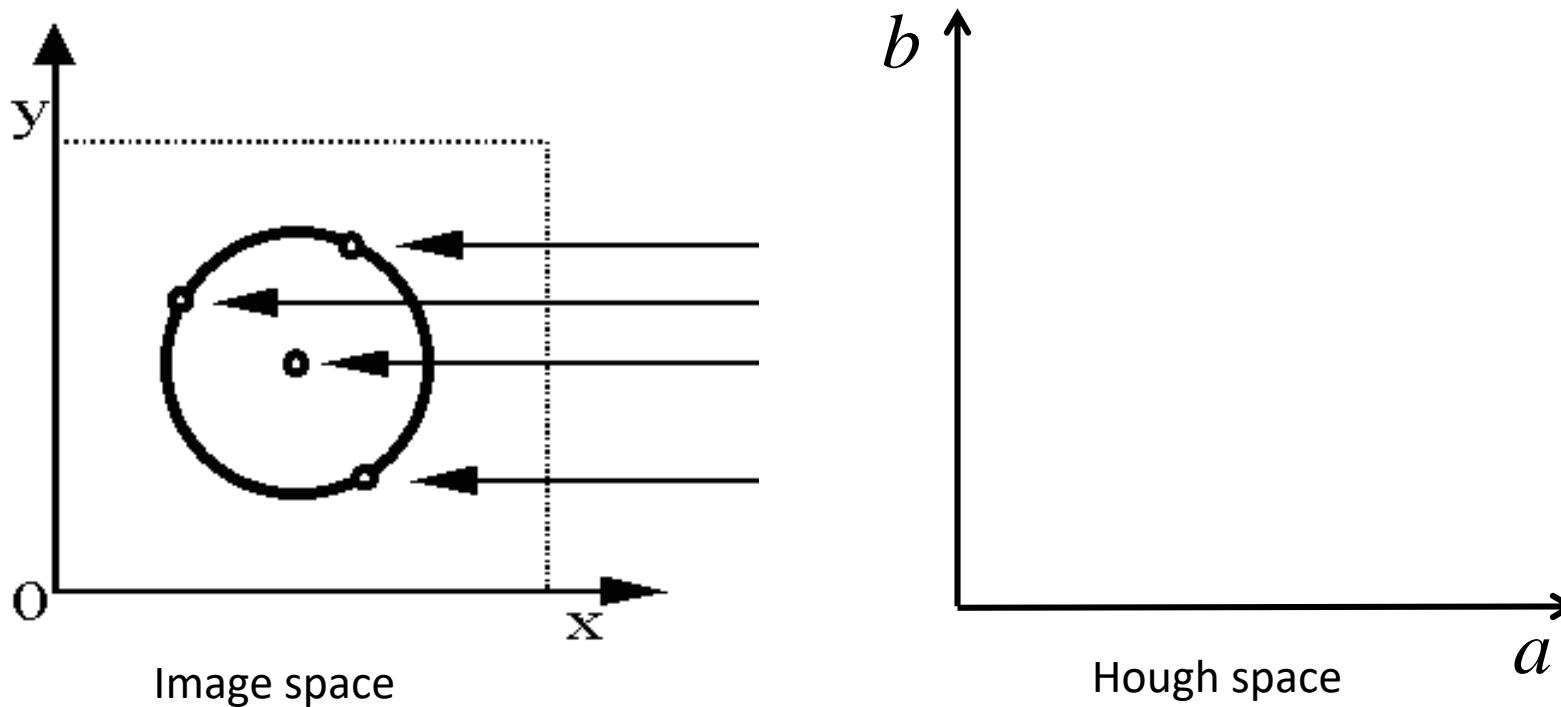
- These extensions can be applied in general:
line, circles, squares, general shapes...

Hough transform for circles

- Circle parameters: center (a, b) and radius r

$$(x_i - a)^2 + (y_i - b)^2 = r^2$$

- Example of center detection at **known radius r**

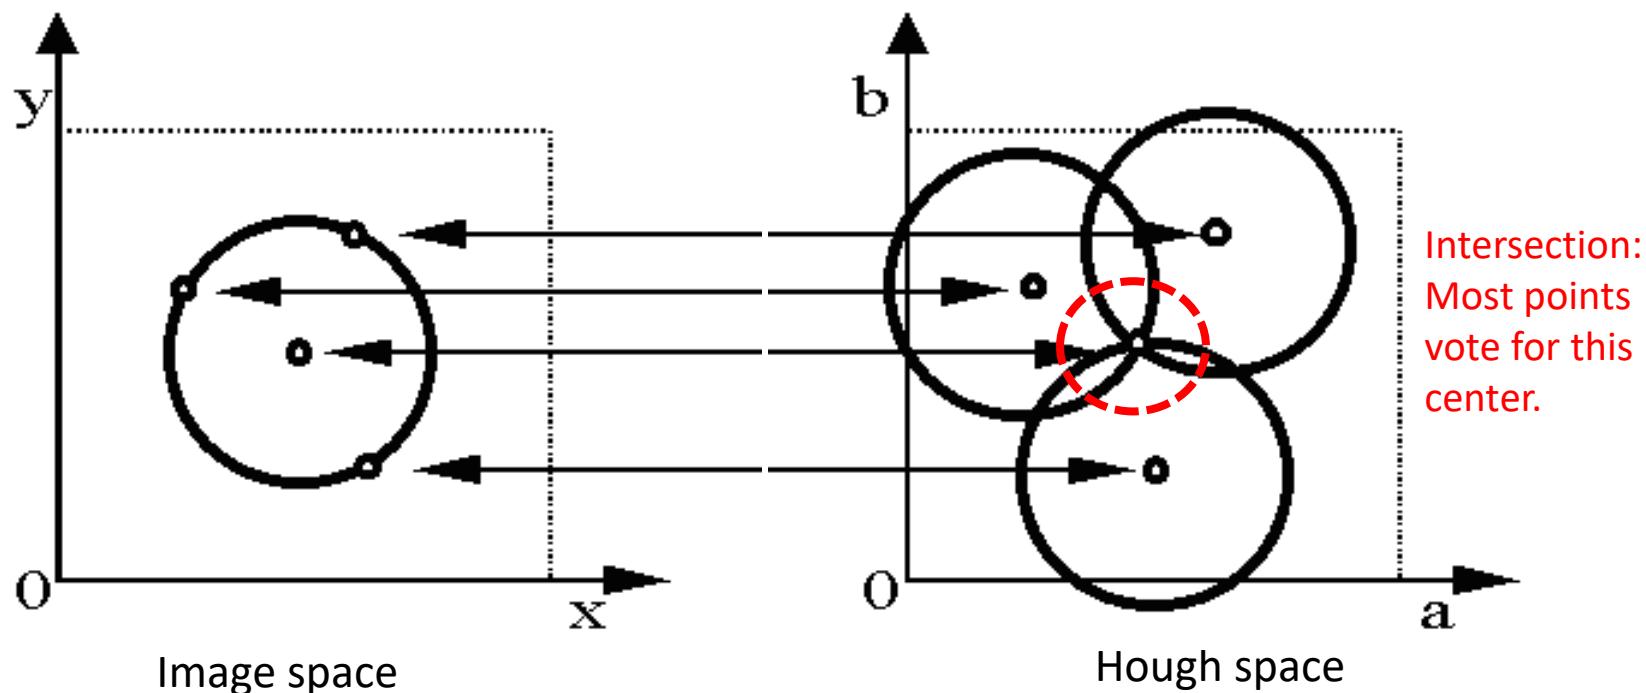


Hough transform for circles

- Circle parameters: center (a, b) and radius r

$$(x_i - a)^2 + (y_i - b)^2 = r^2$$

- Example of center detection at known radius r

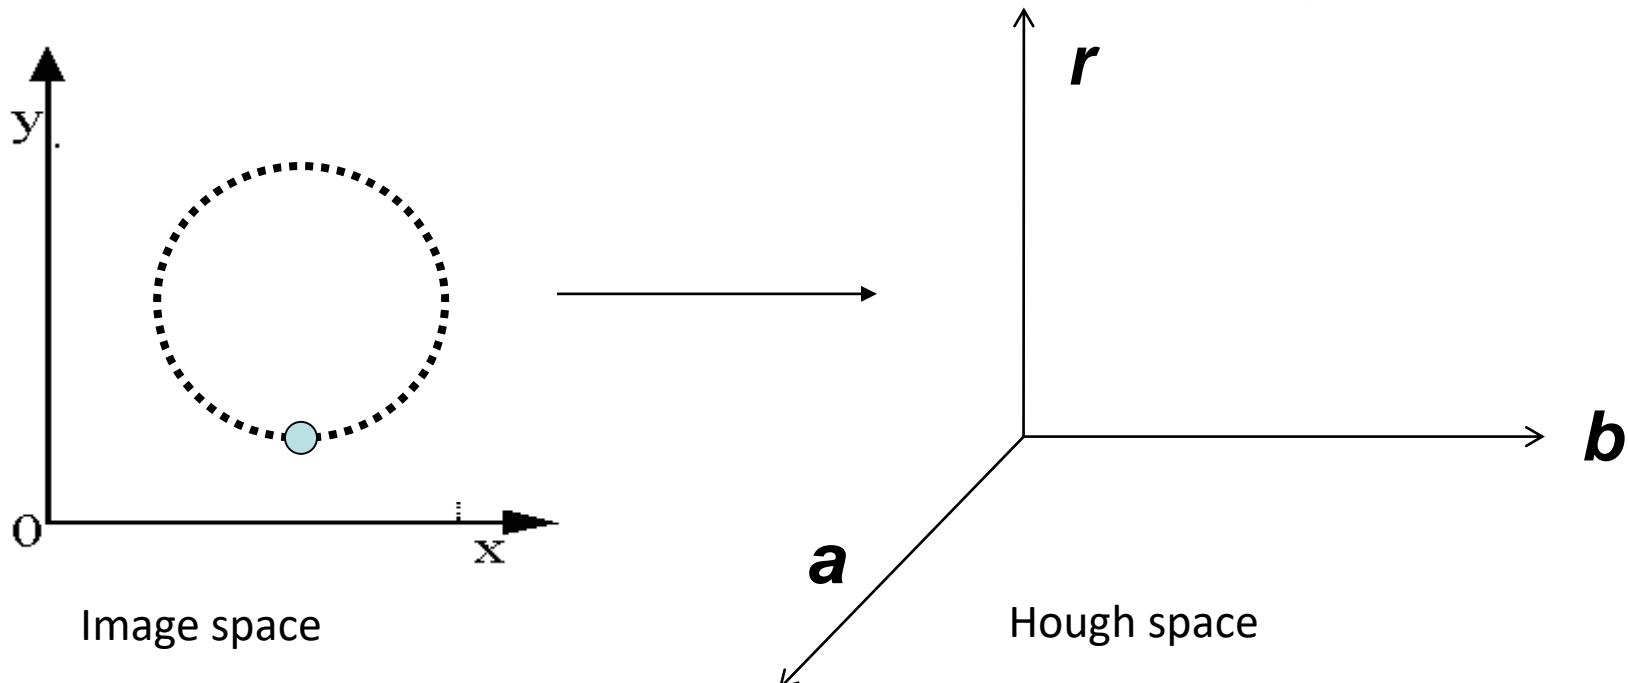


Hough transform for circles

- Circle parameters: center (a, b) and radius r

$$(x_i - a)^2 + (y_i - b)^2 = r^2$$

- Unknown radius r – How many dimensions in Hough Space?

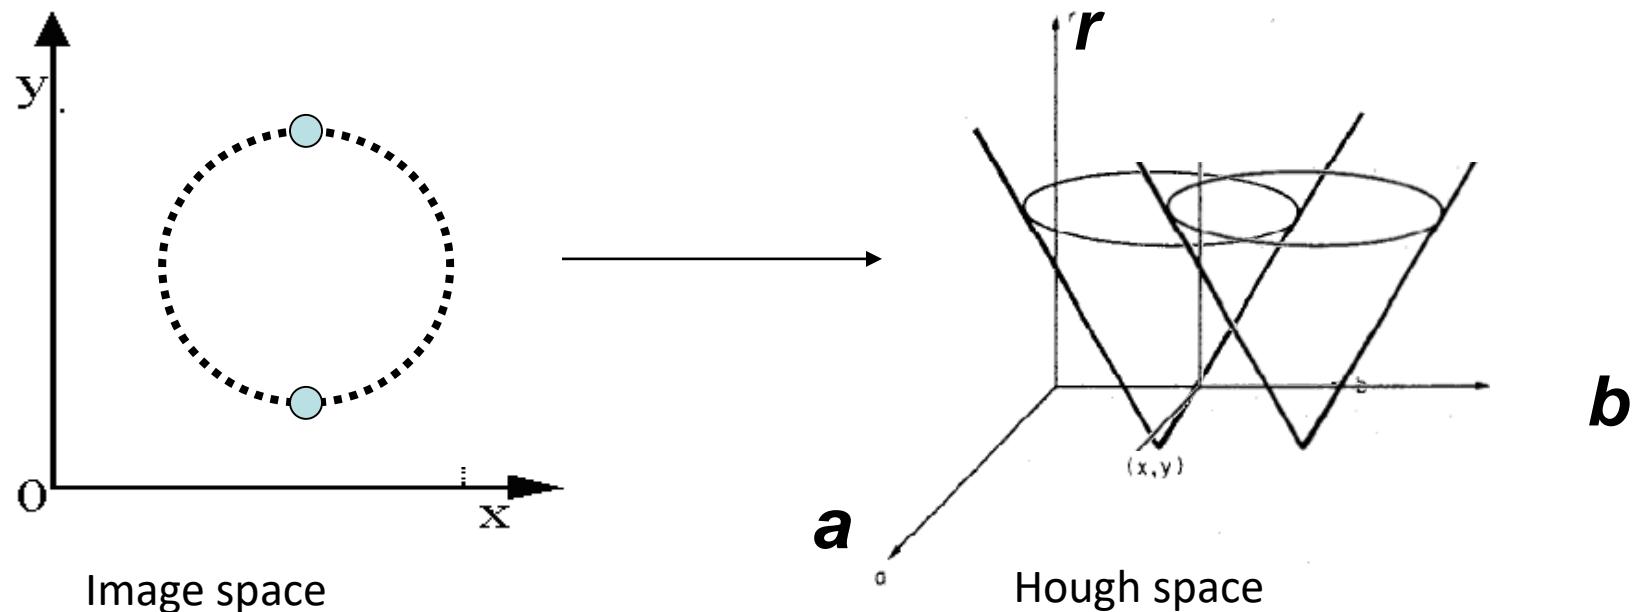


Hough transform for circles

- Circle parameters: center (a, b) and radius r

$$(x_i - a)^2 + (y_i - b)^2 = r^2$$

- Unknown radius r



Hough transform for circles

- Circle parameters: center (a, b) and radius r

$$(x_i - a)^2 + (y_i - b)^2 = r^2$$

- Unknown radius r

But assume we know the gradient direction!

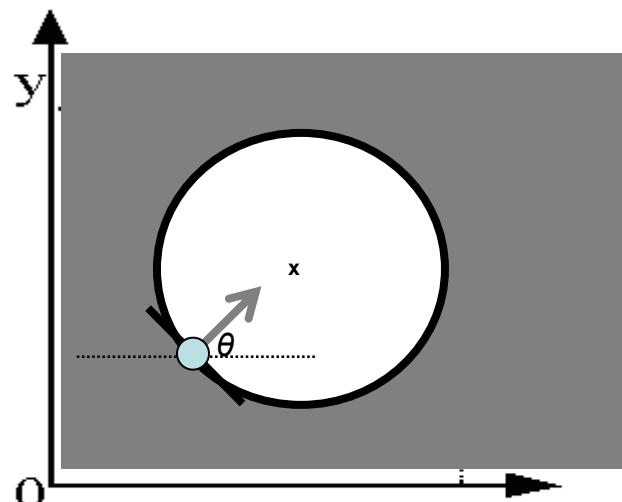
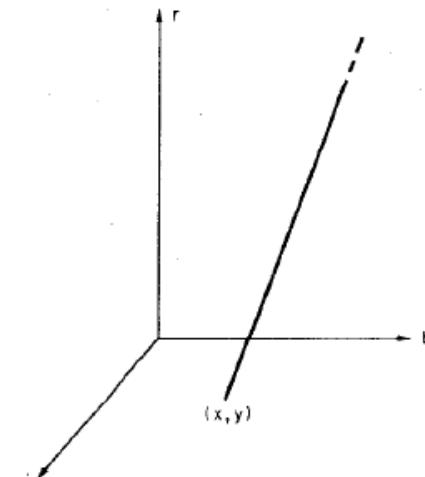


Image space



Hough space

Hough transform for circles

For each edge pixel (x, y) :

For each radius value r :

For each gradient direction θ :

// or use the estimated direction only

$$a = x - r \cos(\theta)$$

$$b = y + r \sin(\theta)$$

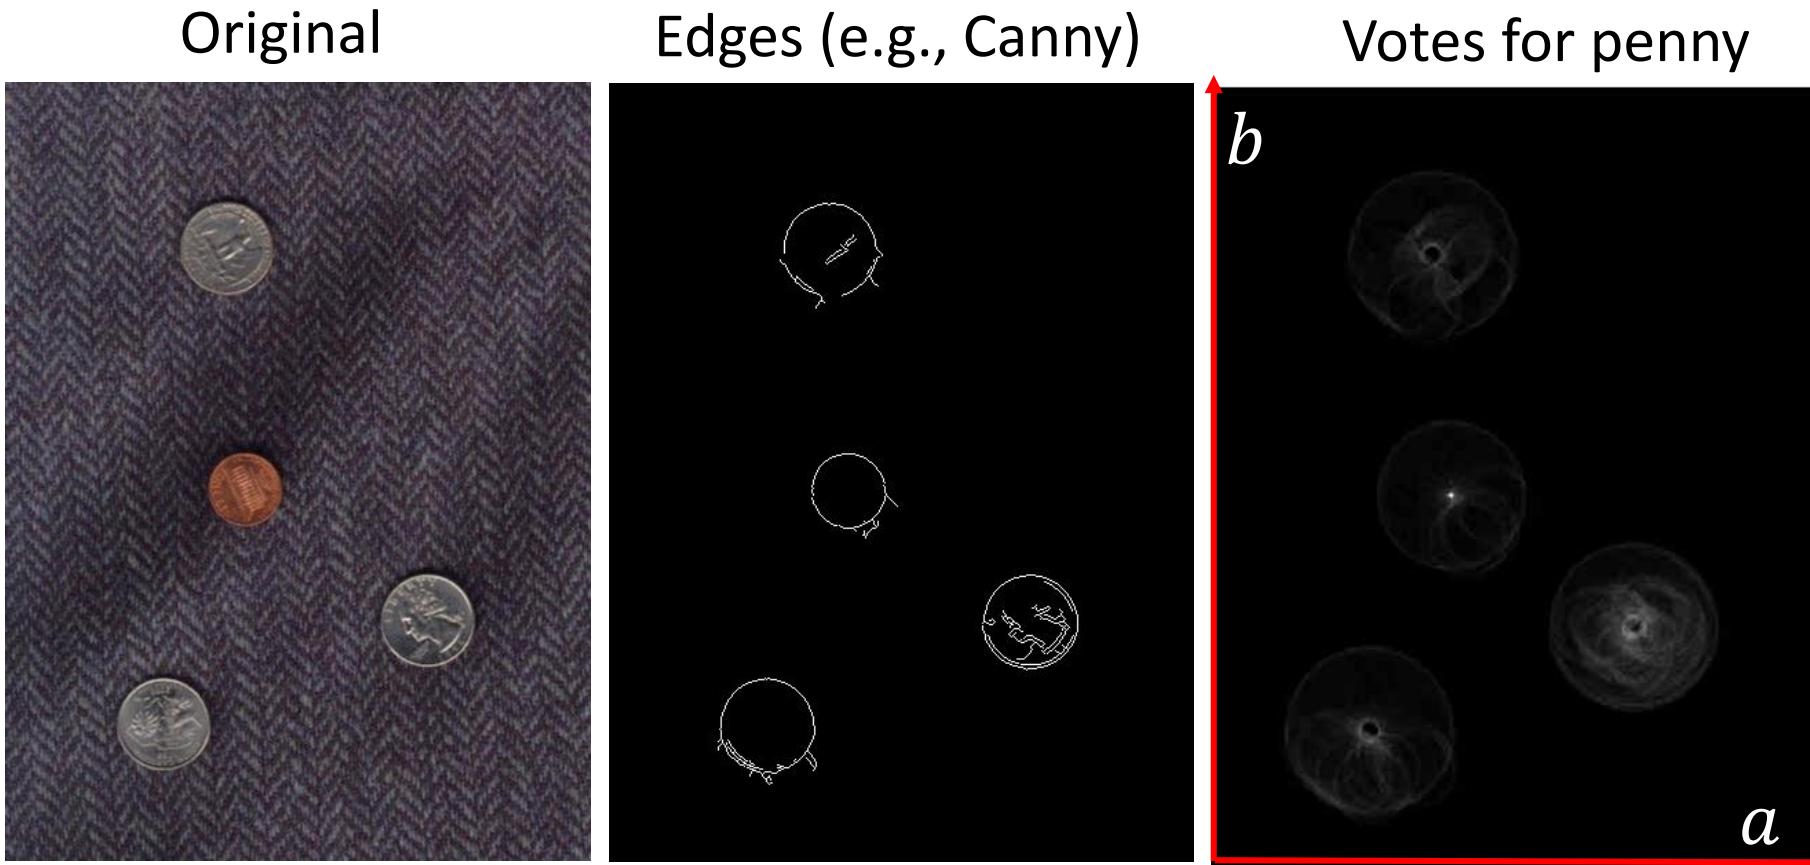
$$H[a, b, r] += 1$$

end for

end for

end for

Hough circle detection in action!



$$(x_i - a)^2 + (y_i - b)^2 = r^2$$

Given a known radius:

$$r = 50 \text{ pixels},$$

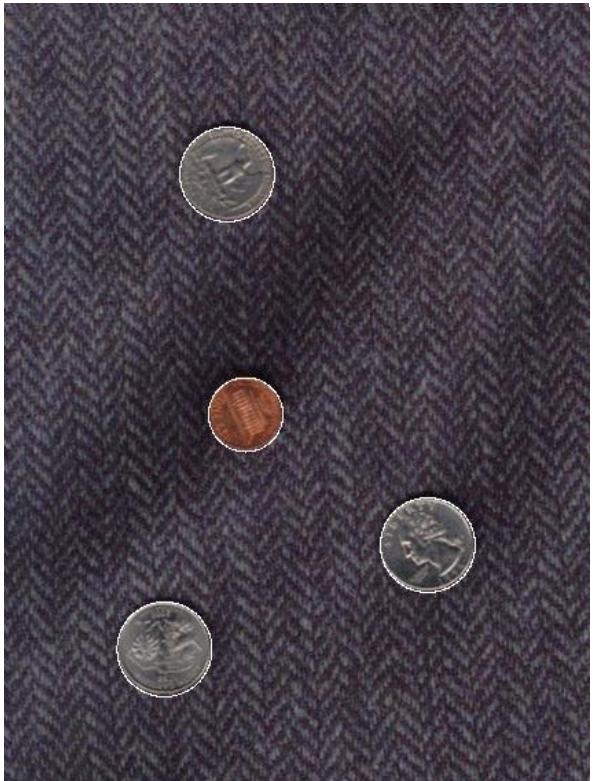
what are the center coordinates, i.e.,

$$a =? \text{ pixels}, \\ b =? \text{ pixels}.$$

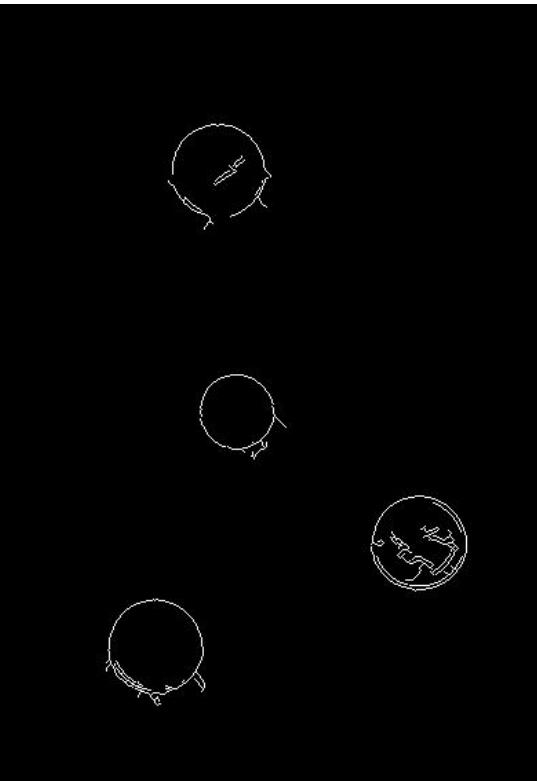
Comment: here we use a separate HT for each coin size.

Hough circle detection in action!

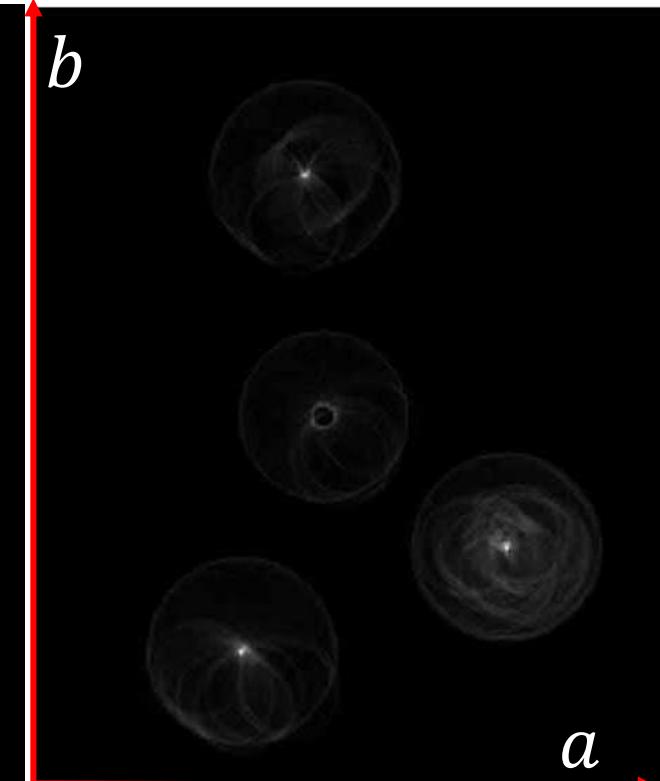
Combined detections



Edges (e.g., Canny)



Votes for 25cent



$$(x_i - a)^2 + (y_i - b)^2 = r^2$$

Given a known radius:

$$r = 50 \text{ pixels},$$

what are the center coordinates, i.e.,

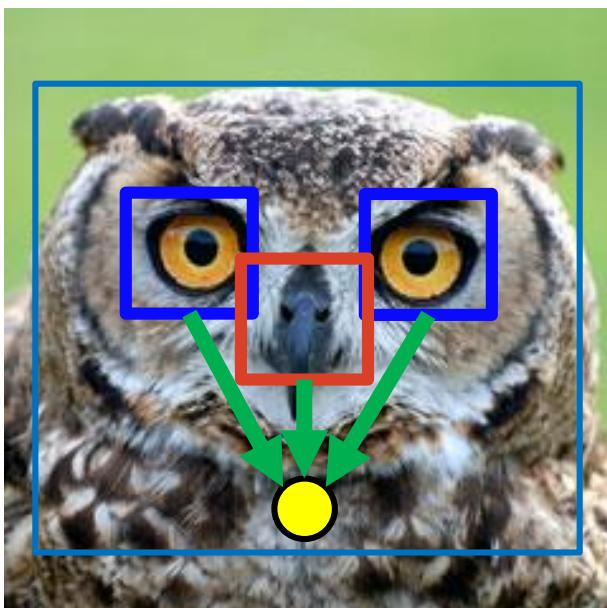
$$a = ? \text{ pixels}, \\ b = ? \text{ pixels}.$$

Comment: here we use a separate HT for each coin size.

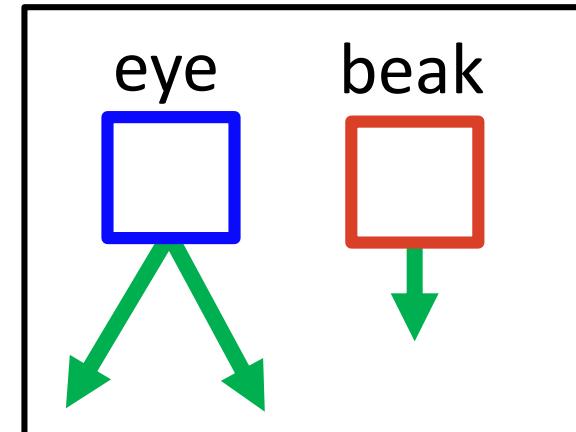
Generalized Hough transform (GHT)

Building a model to detect objects by GHT – intuition:

- Assume we know how to **detect parts** (recognize+localize), i.e., eyes and beak of an owl. **Task:** *create an owl head detector.*
- **Encode parts** by displacements to the neck center.

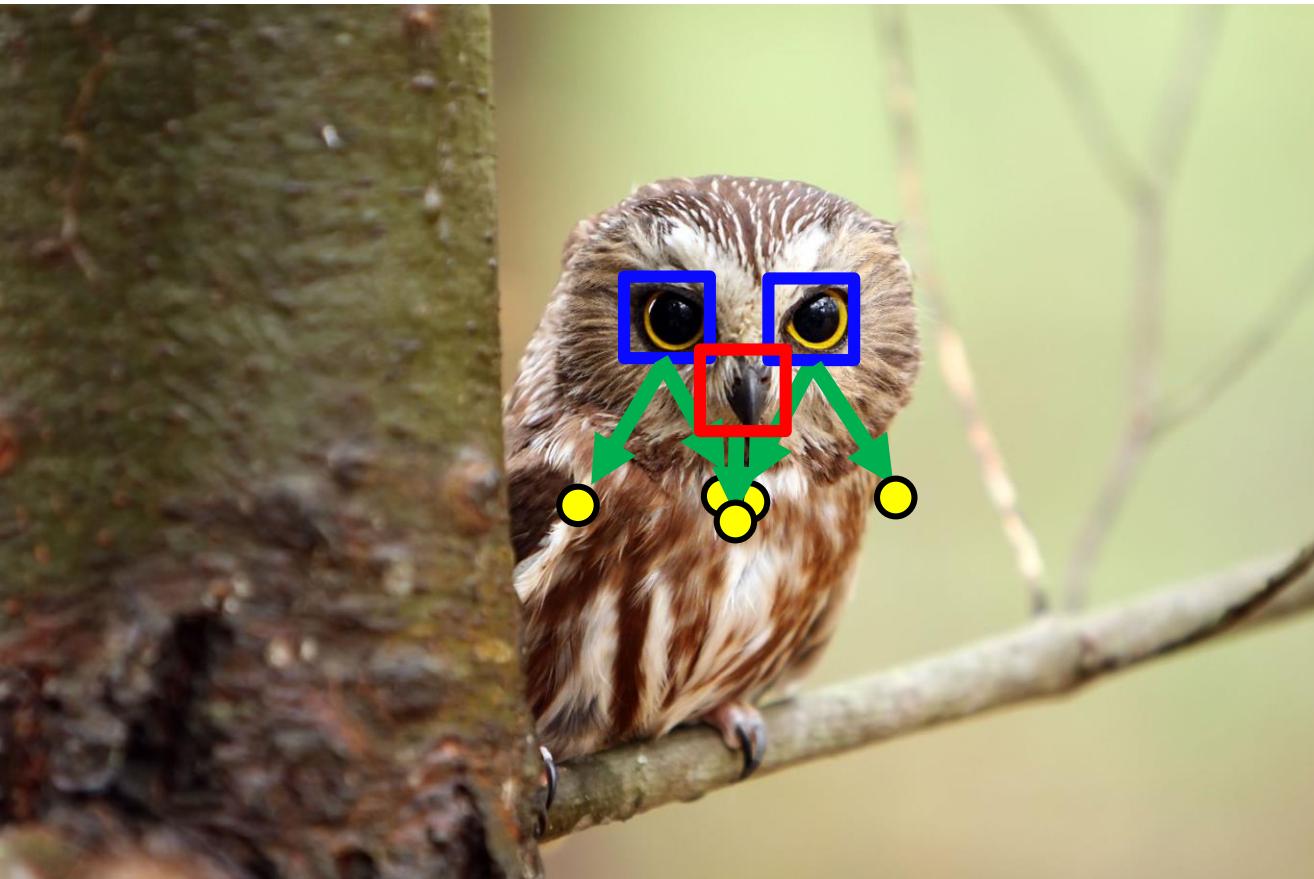


The owl head model:
Given a part, where is the neck center?

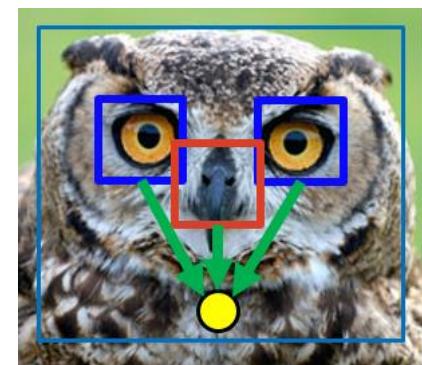
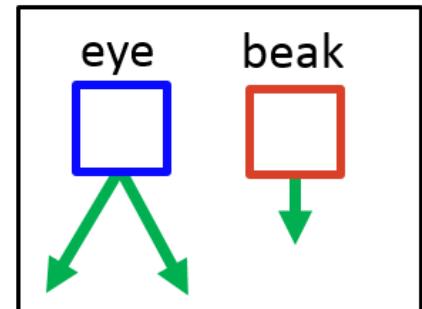


Generalized Hough transform (GHT)

- Detection – intuition

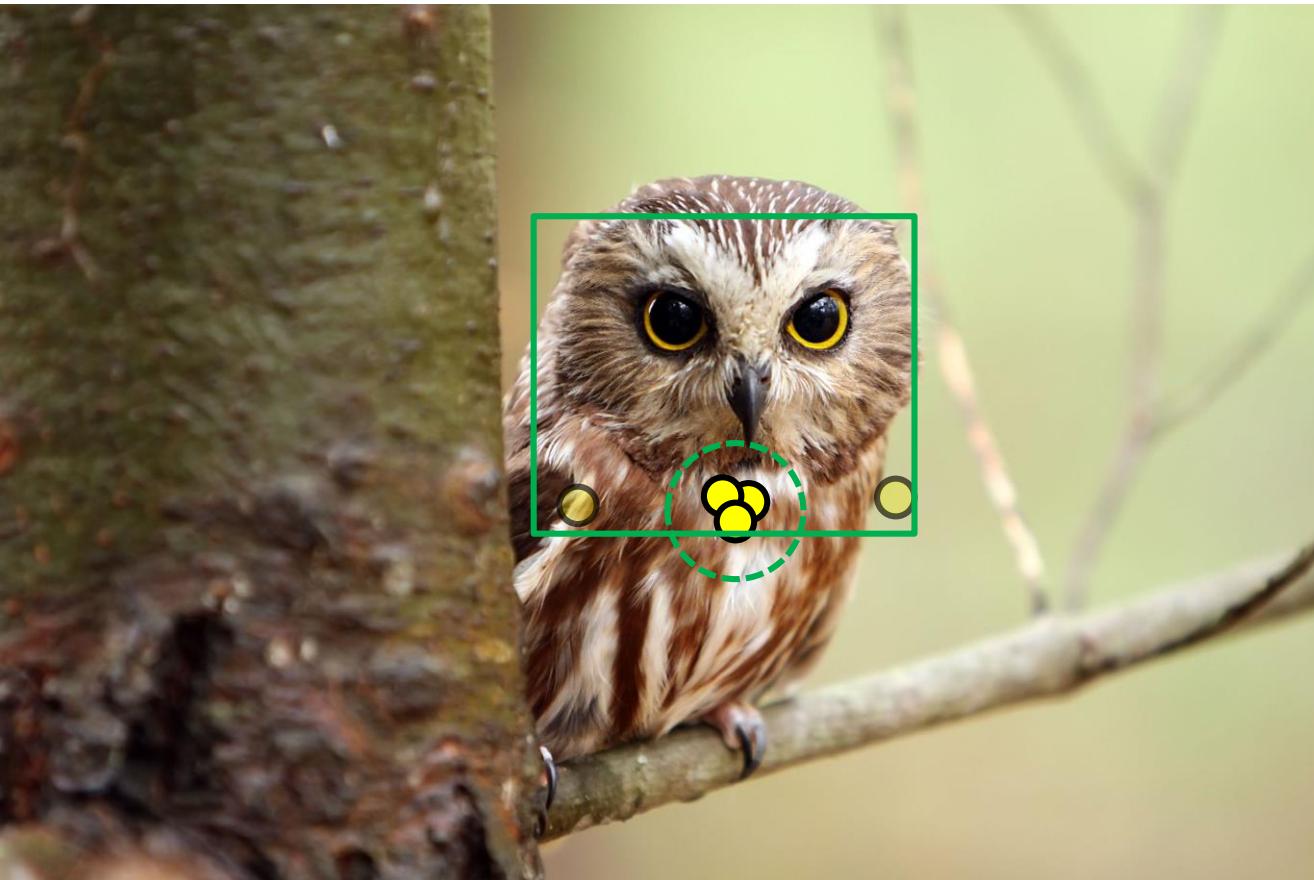


The owl head model:

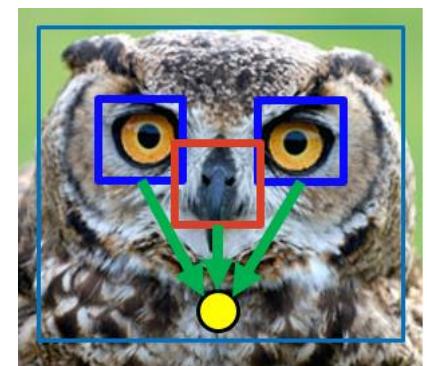
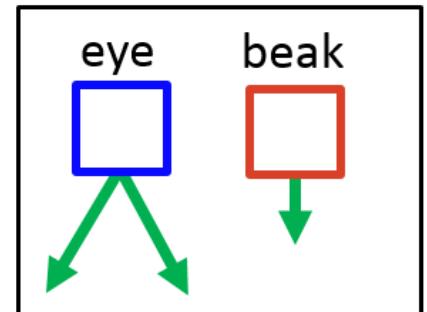


Generalized Hough transform (GHT)

- Detection – intuition

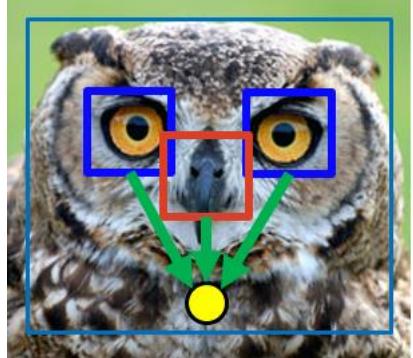


The owl head model:

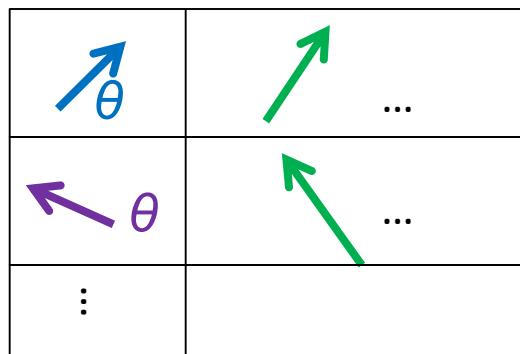
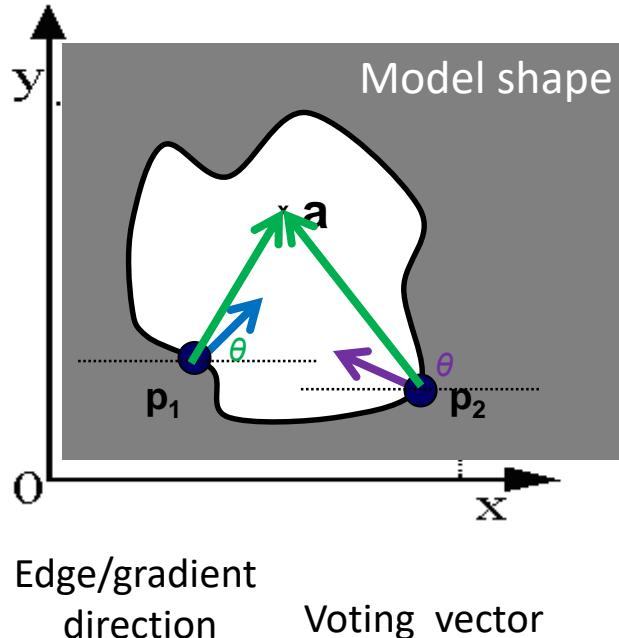
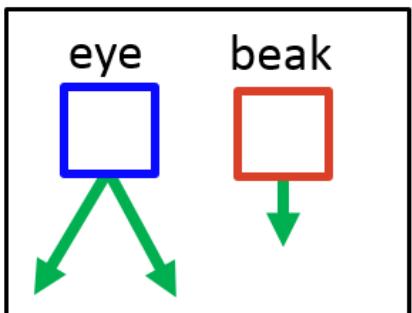


GHT for shape-based models

- Define the shape **model** by edge points and a **reference point**.



The owl head model:



Model learning:

For each edge point calculate the displacement vector to the reference point:

$$\mathbf{r} = \mathbf{a} - \mathbf{p}_i.$$

Collect displacements in **table**, indexed by gradient **direction** θ .

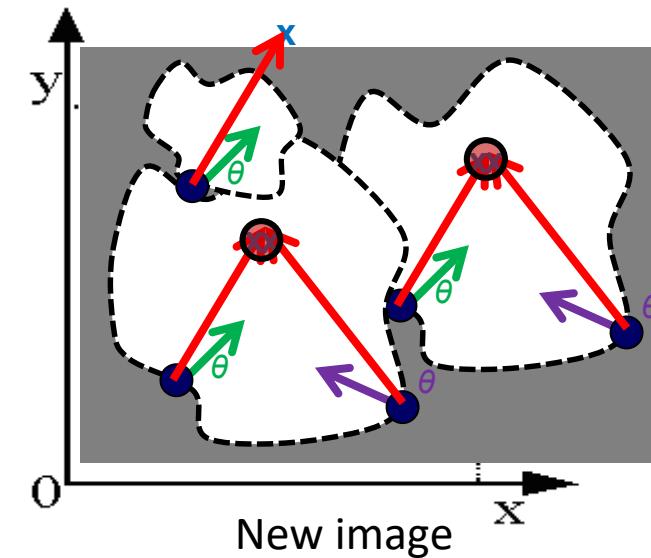
GHT for shape-based models

Detection:

For each edge point:

- Use its gradient **orientation** θ to **index** into the table.
- Use the **displacement vectors** r to cast a **vote** for the **center**.

$\nearrow \theta$	$\nearrow \theta$
$\nwarrow \theta$	$\nearrow \theta$
:	



Assumption: the only transformation is the translation (orientation+scaling are fixed)

Voting: Practical advices

- First minimize irrelevant responses
(use only edges with significant magnitude of gradient)
- Appropriately discretize the parametric space
 - Too coarse: votes from different lines fall into the same accumulator
 - Too fine: losing lines – due to noise, collinear points cast votes into nearby (BUT DIFFERENT) accumulators.
- Vote for neighboring cells as well (smoothing accumulator by Gaussian filtering)
- Use the gradient direction to reduce dof by 1.

Hough transform: +/-

Pros

- Each point is processed **independently**:
 - **robustness** to partial **occlusion**,
 - **highly parallelizable**.
- **Robustness to noise**: noise will unlikely contribute consistently to a single cell
- Can **detect multiple instances** of a single model in one pass.

Cons

- Time **complexity increases exponentially** with the number of free parameters.
- Spurious shapes may generate **false local maxima** in the parametric space.
- Quantization: **Not particularly easy** to choose a **proper** accumulator **cell size** – Application dependent!

References

- David A. Forsyth, Jean Ponce, Computer Vision: A Modern Approach (2nd Edition),
(prva izdaja dostopna na spletu)
- R. Szeliski, Computer Vision: Algorithms and Applications, 2010
- R. Hartley, A. Zisserman, Multiple View Geometry in Computer Vision,
2nd Edition, Cambridge University Press, 2004
- Kristen Grauman, „Computer Vision“, lectures